

Basic Image Processing Algorithms

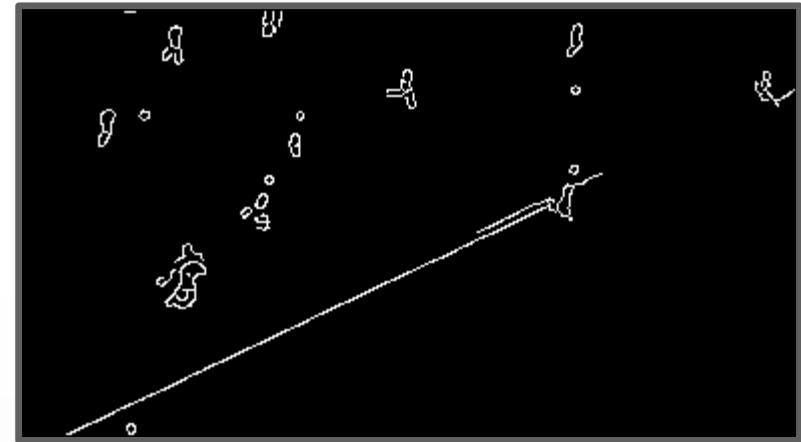
PPKE-ITK

Lecture 3.

Line detection with Hough Transfrom

Hough Transformation

- An example of Canny edge detector...



- ...where straight lines are not detected perfectly.
- The objective of the Hough transformation is to find the lines on a binary image, from fragments/points of the line.

Finding lines in an image

- Option 1:

- Search for the line at every possible position/orientation
- What is the cost of this operation?

- Option 2:

- Use a voting scheme: Hough transform

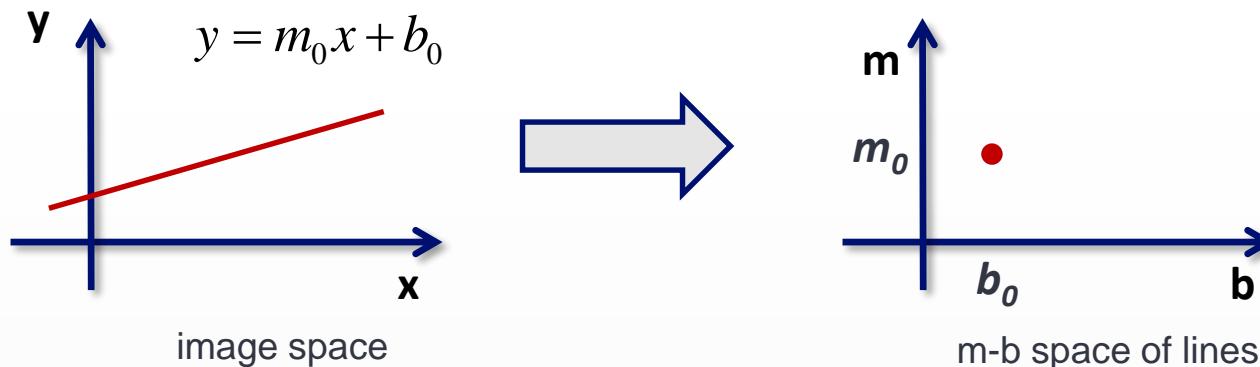
Finding lines in an image

- The basic idea:

- A line can be written in the following form:

$$y = mx + b$$

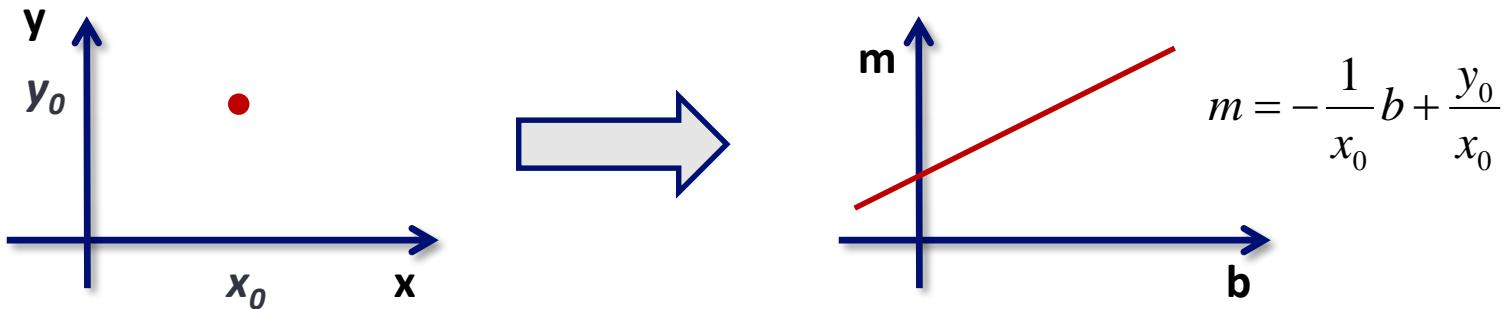
where m is the slope of the line and b is the y-intercept.



- Connection between image (x,y) and the (m,b) spaces

- A **line in the image** corresponds to a **point in "m-b" space**
 - To go from image space to (m-b) space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$

Finding lines in an image



◎ Connection between image (x, y) and (m, b) spaces

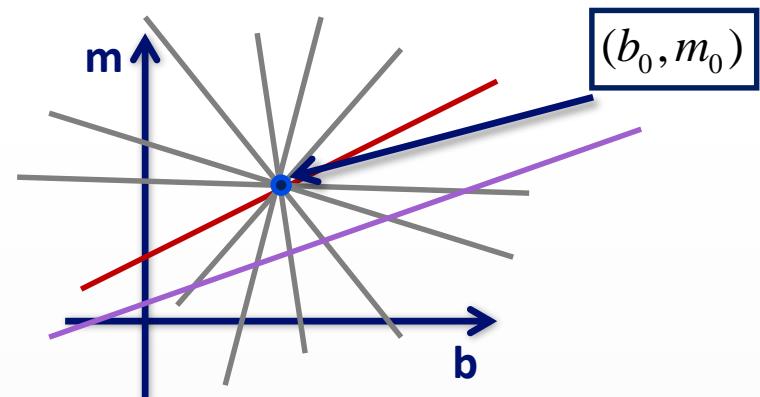
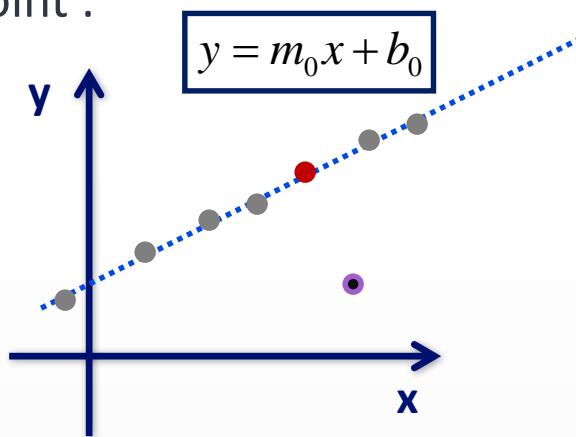
- What does a point (x_0, y_0) in the image space map to?
 - For a fixed $y = y_0$, $x = x_0$ point in the image space, we get a line in the (m, b) space with a slope $-1/x_0$ and an m -intercept: y_0/x_0 :

$$m = -\frac{1}{x_0}b + \frac{y_0}{x_0}$$

Hough Transformation

- The basic idea:

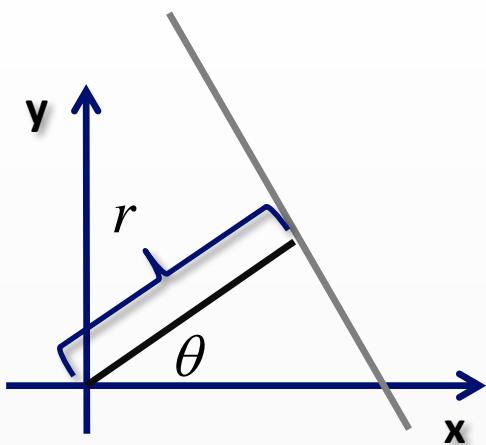
- For the points that lie on the same line in the Euclidian space, their corresponding line in the parameter space will cross each other in one point :



- This point will be $m=m_0$ and $b=b_0$, the slope and intercept of the line in the image space. \rightarrow We have the equation of the line!
- But, there is a problem with this equation of the line: **vertical lines cannot be described** (their slope would be infinite).

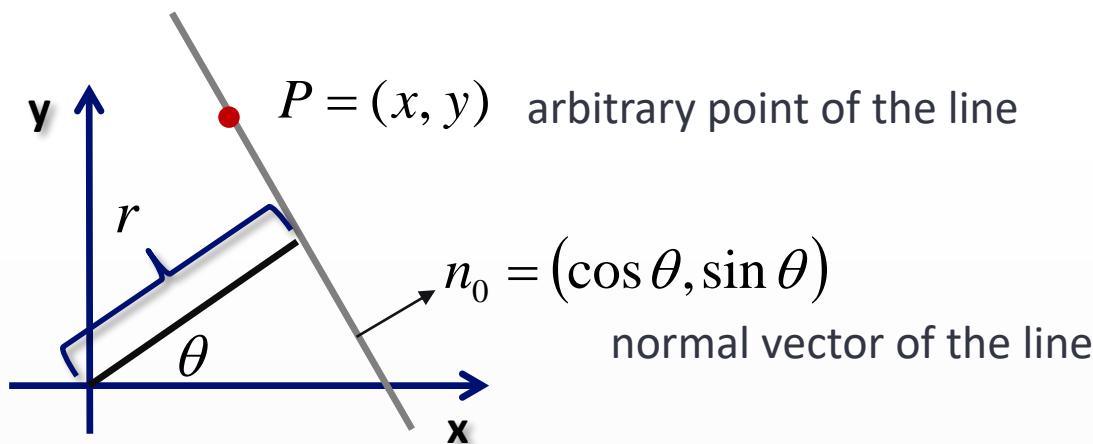
Hough Transformation

- ◎ To be able to describe all possible lines with two scalar parameters, we will use a **polar representation** of the line
- ◎ Each line is described by (r, θ) instead of (m, b) , where
 - r is the perpendicular distance from the line to the origin
 - θ is the angle this perpendicular makes with the x axis



Hough Transformation

- Mathematical basis for using the **polar equation** of the line is the Hesse normal form*: $0 = P \cdot n_0 - r$

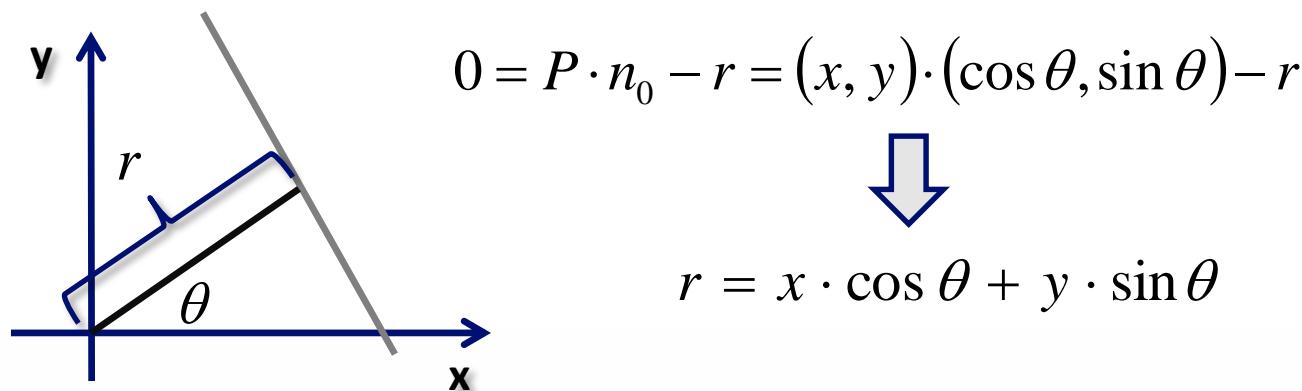


- r : perpendicular distance from the line to the origin
- θ : the angle this perpendicular makes with the x axis

* https://en.wikipedia.org/wiki/Hesse_normal_form

Hough Transformation

- Hesse normal form based polar equation of the line:

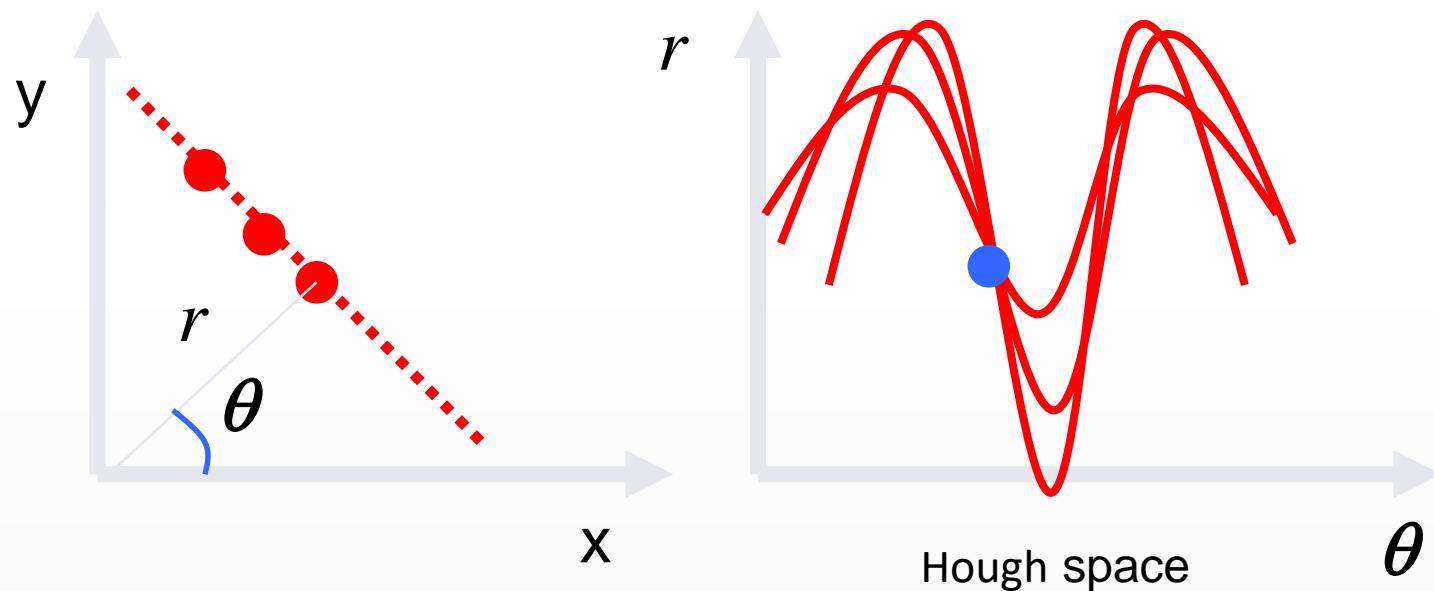


- The (r, θ) parameter space is called Hough space.
- A point in the Euclidian space is a sinusoid in the Hough space, described by the following equation:

$$r(\theta) = x \cdot \cos \theta + y \cdot \sin \theta$$

Hough Transformation

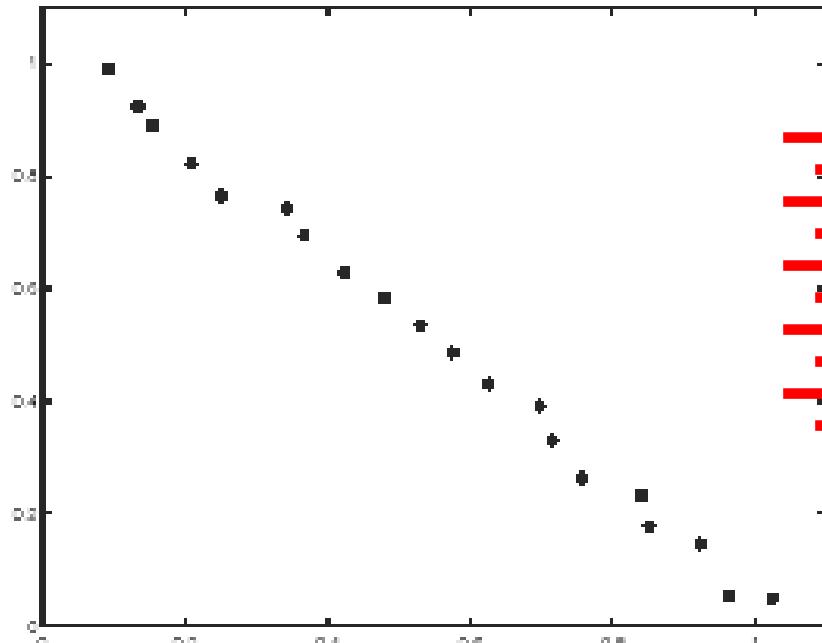
- All the sinusoid curves of the points in one line in the Euclidian space, cross each other in one point in the Hough space.



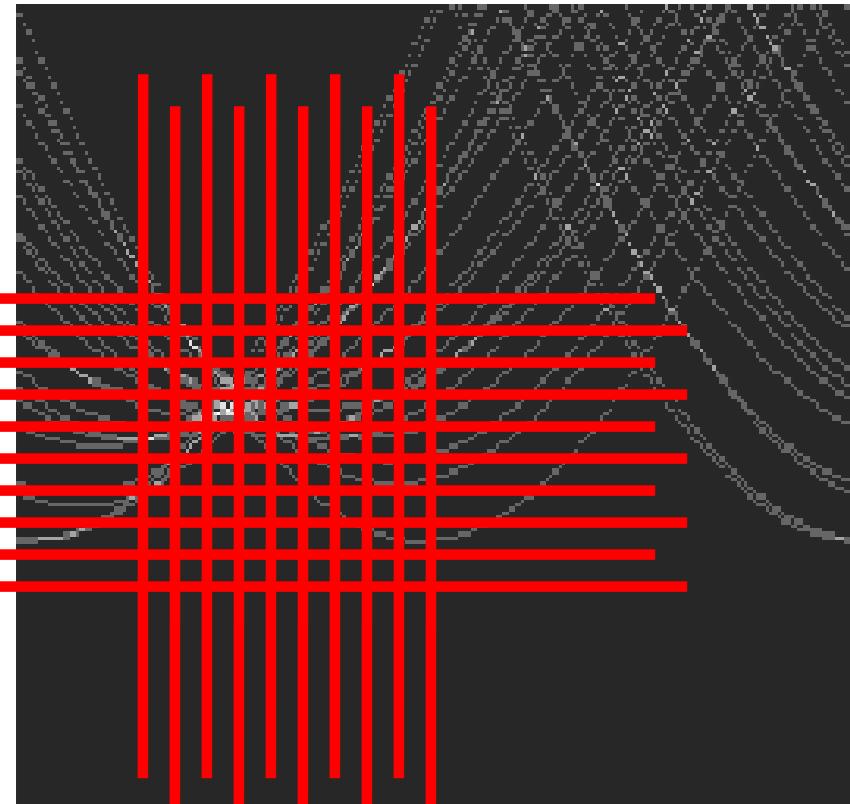
$$r(\theta) = x \cdot \cos \theta + y \cdot \sin \theta$$

Hough Transformation

Noisy data



features



votes

Issue: Grid size needs to be adjusted...

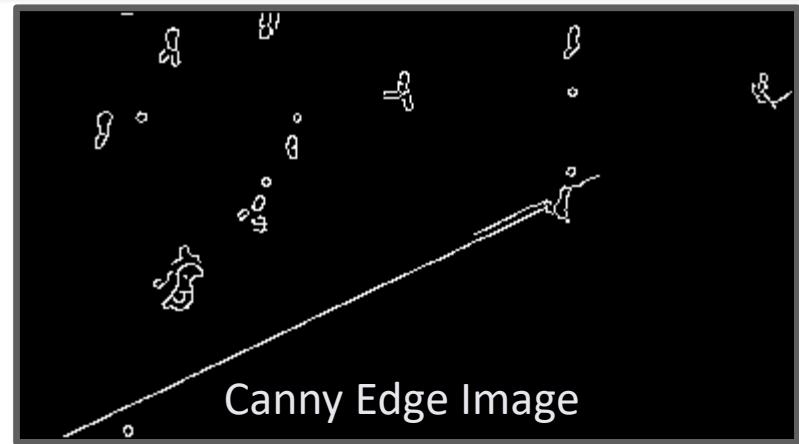
Hough transform algorithm

- Basic Hough transform algorithm
 - 1. for all r, θ : initialize $H[r, \theta] = 0$
 - 2. for each edge point $I[x, y]$ in the image
 - for $\theta = 0$ to 180
$$r = x \cdot \cos \theta + y \cdot \sin \theta$$
$$H[r, \theta] += 1$$
 - 3. Find the value(s) of (r, θ) where $H[r, \theta]$ is maximum
 - 4. The detected line in the image is given by
$$r = x \cdot \cos \theta + y \cdot \sin \theta$$
- What's the running time (measured in # votes)?

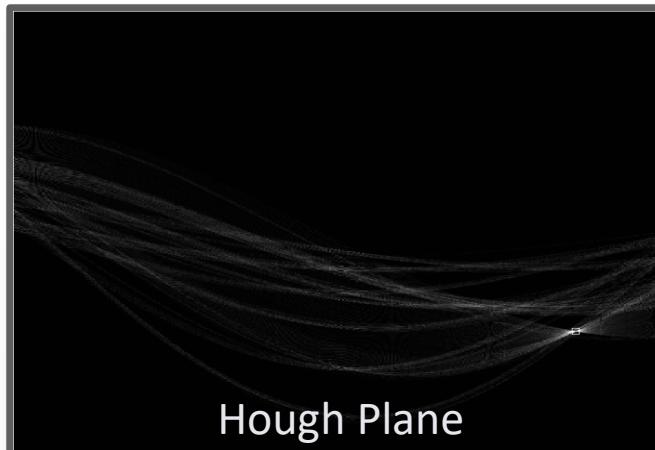
Hough Transformation



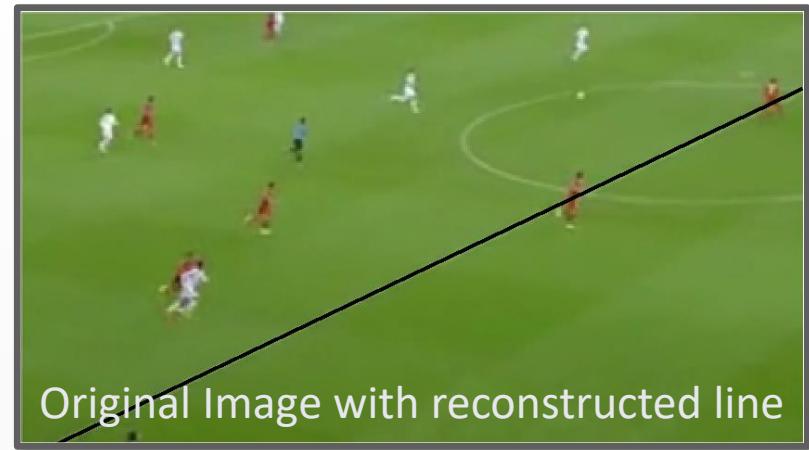
Original Image



Canny Edge Image



Hough Plane



Original Image with reconstructed line

Extensions

- Extension 1: Use the image gradient
 1. same
 2. for each edge point $I[x,y]$ in the image
 - compute unique (r, θ) based on local image gradient at (x,y)
 - $H[r, \theta] += 1$
 3. same
 4. same
- Extension 2
 - give more votes for stronger edges
- Extension 3
 - change the sampling of (r, θ) to give more/less resolution
- Extension 4
 - The same procedure can be used with circles, squares, or any other shape

Hough demos

- Lines, circles and ellipses:

<http://dersmon.github.io/HoughTransformationDemo/>

- Circle : <http://www.markschulze.net/java/hough/>

Image Enhancement

What is Image Enhancement?

- **Image enhancement** is the manipulation or transformation of the image to improve the visual appearance or to help further automatic processing steps.
- There is no general theory behind it, the result is highly application dependent and subjective.
 - e.g. in many cases the goal is to improve the quality for human viewing (Medical Imaging, Satellite Images)
- Enhancement is closely related to image recovery.
- Examples:
 - Contrast enhancement
 - Edge enhancement
 - Noise removal/smoothing

Types of Image Enhancement

- There are two main categories:
 - Spatial Domain Methods
 - Frequency Domain Methods
- In the Spatial Domain we are directly manipulating pixel values, through..
 - Point-wise Intensity Transformation
 - Histogram Transformations
 - Spatial Filtering
 - LSI (*Linear Shift-Invariant*)
 - Non-Linear
 - etc.

The Histogram of an Image

- **Histogram:**

$h(k)$ = the number of pixels on the image with value k .



Original Image*

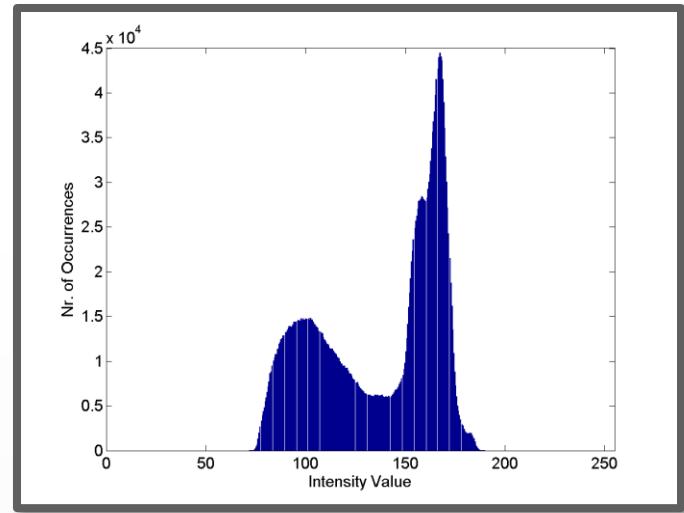


Image Histogram

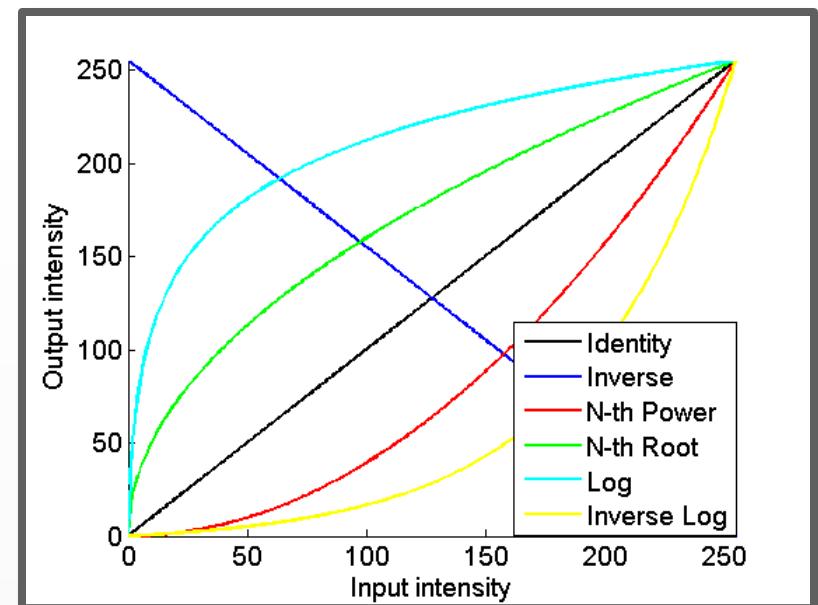
- The histogram normalized with the total number of pixels gives us the ***probability density function*** of the intensity values.

* Modified version of Riverscape with Ferry by Salomon van Ruysdael (1639)

Point-wise Intensity Transformation

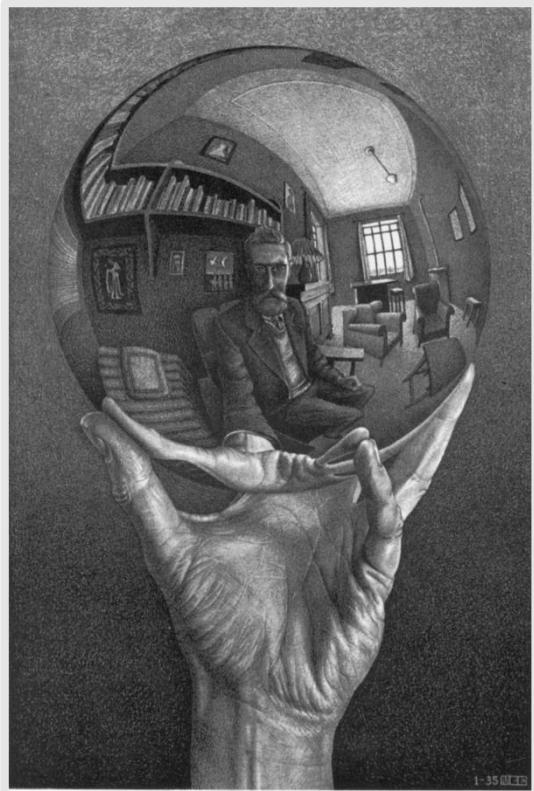
- Point wise transformations are operating directly on pixel values, independently of the values of its neighboring pixels.
- We can describe the transformation as follows:
 - Let x and y be two grayscale images, and let T be a point-wise image enhancement transformation that transforms x to y :

$$y(n_1, n_2) = T[x(n_1, n_2)]$$

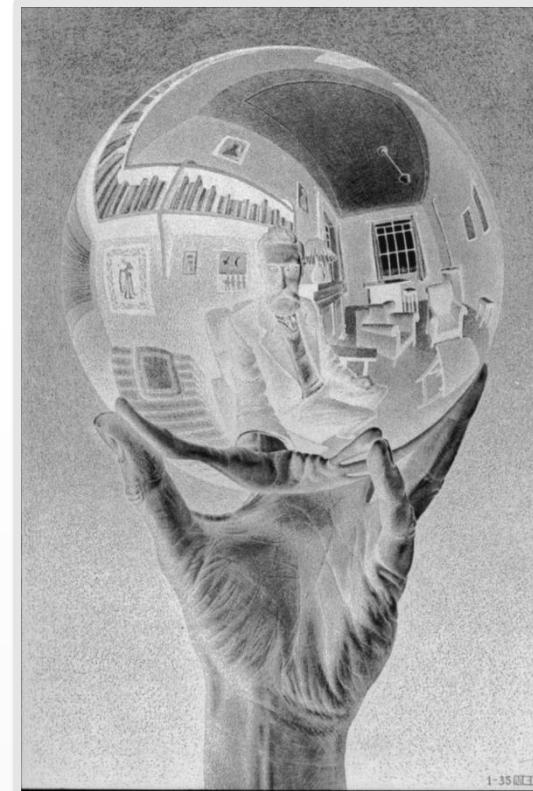


Point-wise Intensity Transformation

- ◎ Inverse transformation: $y(n_1, n_2) = 255 - x(n_1, n_2)$



Original Image*



Inverse Image

*Hand with Reflecting Sphere by M. S. Escher (1935)

Point-wise Intensity Transformation

- ◎ **Log transformation:** $y(n_1, n_2) = c \cdot \log(x(n_1, n_2) + 1)$
 - Expands low and compresses high pixel value range



Original Image*



Log Image



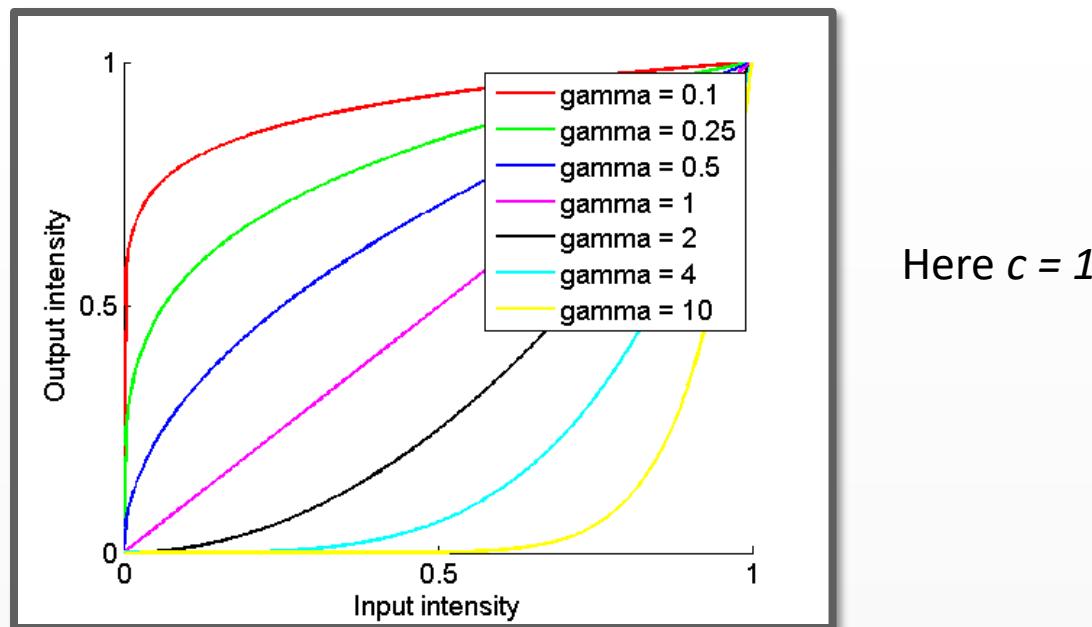
Log Image
after histogram stretching

* Abbaye du Thoronet by Lucien Hervé (1951)

Point-wise Intensity Transformation

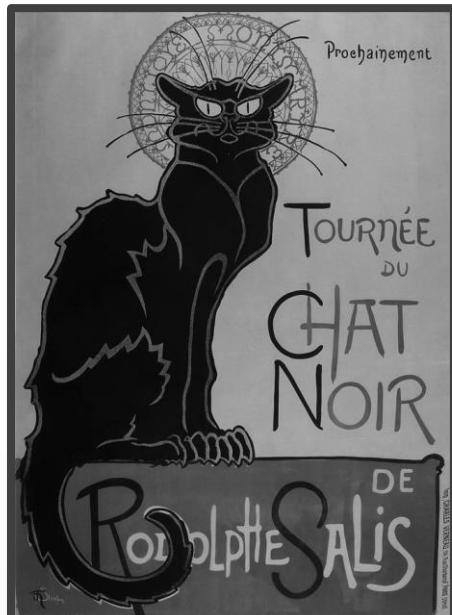
- ◎ **Power-law transformation:** $y(n_1, n_2) = c \cdot x(n_1, n_2)^\gamma$

- Commonly referred to as ***gamma transformation***
- Originally it was developed to compensate the input-output characteristics of CRT displays.
- The expended/compressed region depends on γ :

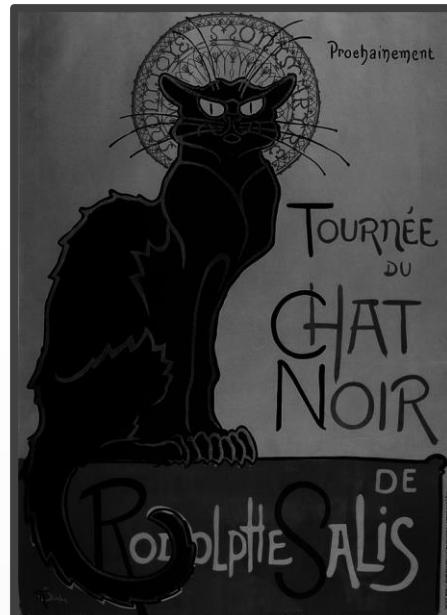


Point-wise Intensity Transformation

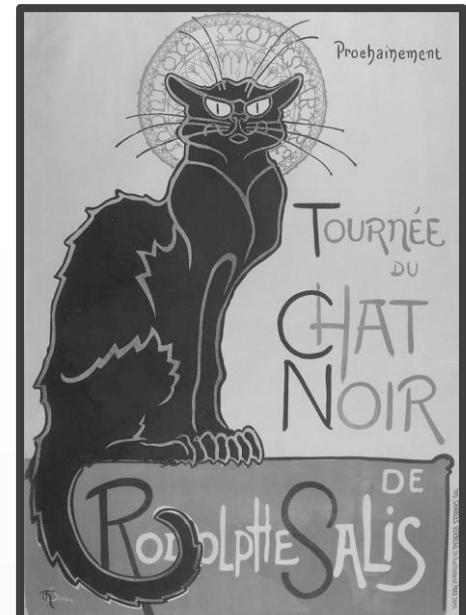
- ◎ Power-law transformation: $y(n_1, n_2) = c \cdot x(n_1, n_2)^\gamma$



Original Image*



$\gamma = 2$

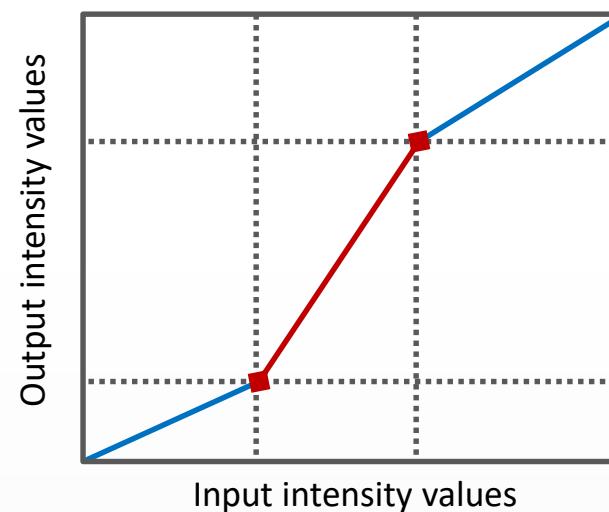


$\gamma = 0.5$

* Le chat Noir, Poster of Théophile Steinlen (1896)

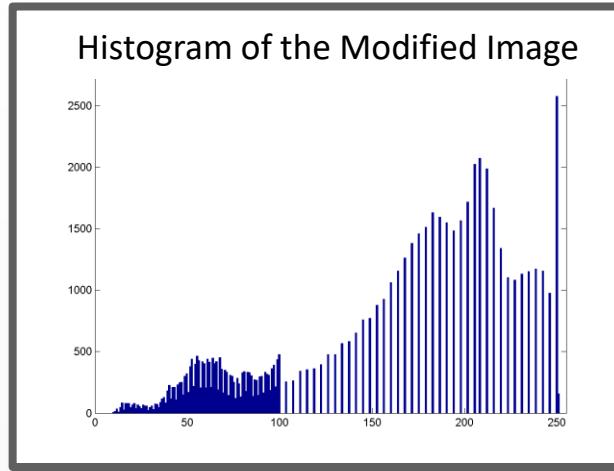
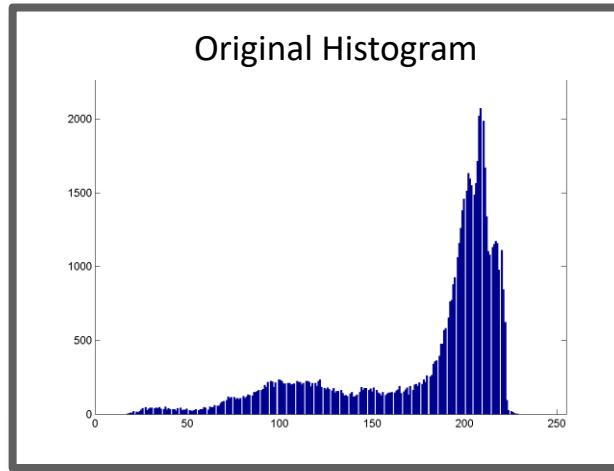
Dynamic Range Expansion

- Piecewise linear expansion/compression of predefined intensity ranges:



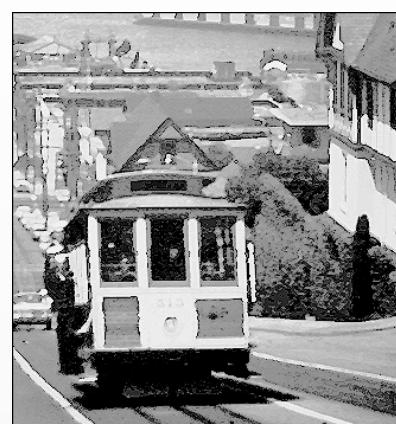
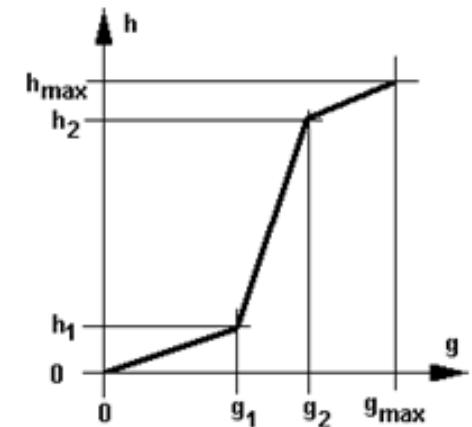
- The red intensity range was expanded, while the blue ranges were compressed.

Dynamic Range Expansion



Dynamic Range Expansion

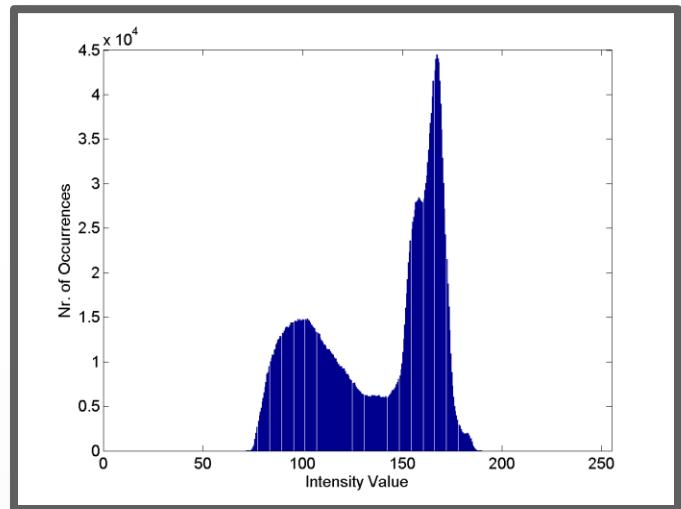
- Example: extracting the intensity values from the $[g_1, g_2]$ interval to a wider $[h_1, h_2]$ domain
 - Enhanced contrast in the selected region, details are better observable and distinguishable.
 - In the remaining image regions the contrast decreases



Histogram Transformations

◎ Histogram Stretching:

- Based on the histogram we can see that the image does not use the whole range of possible intensities:
 - Minimum intensity level: 72
 - Maximum intensity level: 190
- With the following transformation we can stretch the intensity values so they use the whole available range:

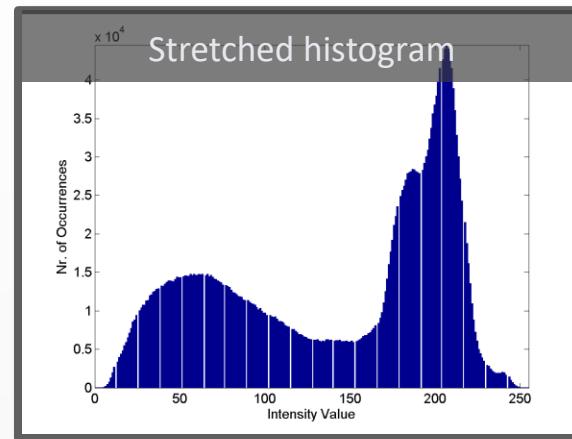
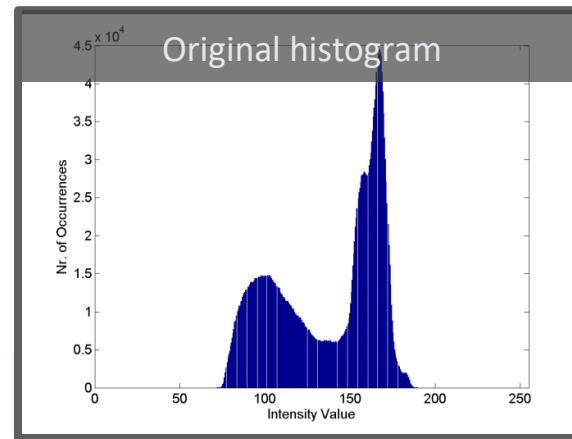


$$y(n_1, n_2) = \frac{255}{x_{\max} - x_{\min}} \cdot (x(n_1, n_2) - x_{\min})$$

$$x_{\max} = \max_{n_1, n_2}(x(n_1, n_2)) \quad x_{\min} = \min_{n_1, n_2}(x(n_1, n_2))$$

Histogram Transformations

○ Histogram Stretching:



Histogram Transformations

- Histogram stretching with various transfer functions:

- Linear:

$$y(n_1, n_2) = \frac{255}{x_{\max} - x_{\min}} \cdot (x(n_1, n_2) - x_{\min}) = 255 \cdot \frac{x(n_1, n_2) - x_{\min}}{x_{\max} - x_{\min}}$$

- Quadratic:

$$y(n_1, n_2) = 255 \cdot \left(\frac{x(n_1, n_2) - x_{\min}}{x_{\max} - x_{\min}} \right)^2$$

- Square root

$$y(n_1, n_2) = 255 \cdot \sqrt{\frac{x(n_1, n_2) - x_{\min}}{x_{\max} - x_{\min}}}$$

Histogram stretching - results



original

linear $f()$



quadratic



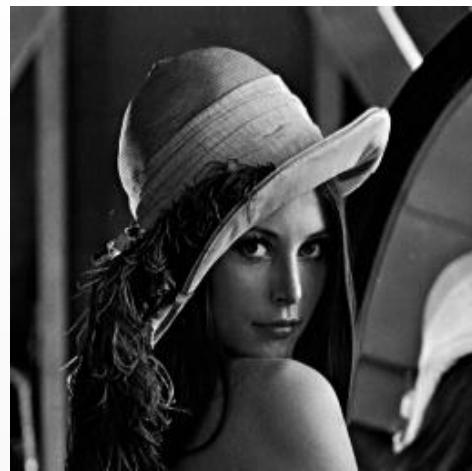
square root

Histogram stretching - results

original



quadratic



linear $f()$

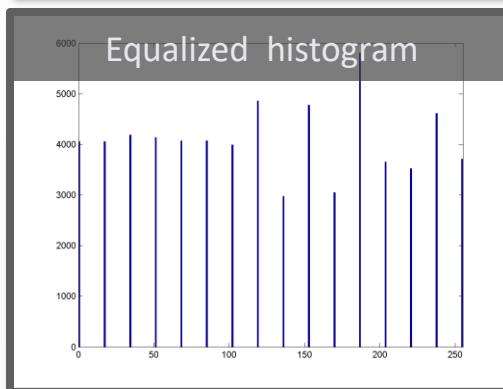
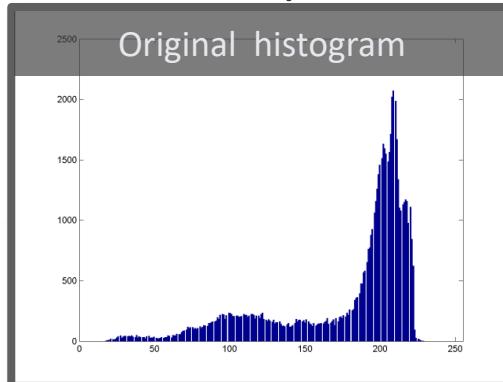
square
root



Histogram Transformations

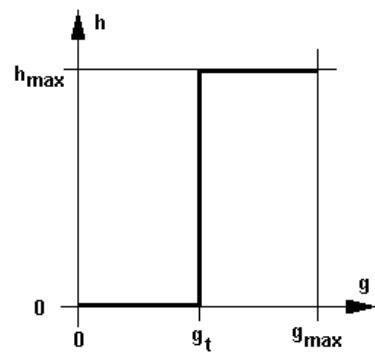
◎ Histogram Equalization:

- The goal is to increase the contrast, by distributing the occurrences of the intensity values evenly through the entire dynamic range.



Histogram equalization background

- Simple thresholding



- For different g_t values



Optimal threshold value

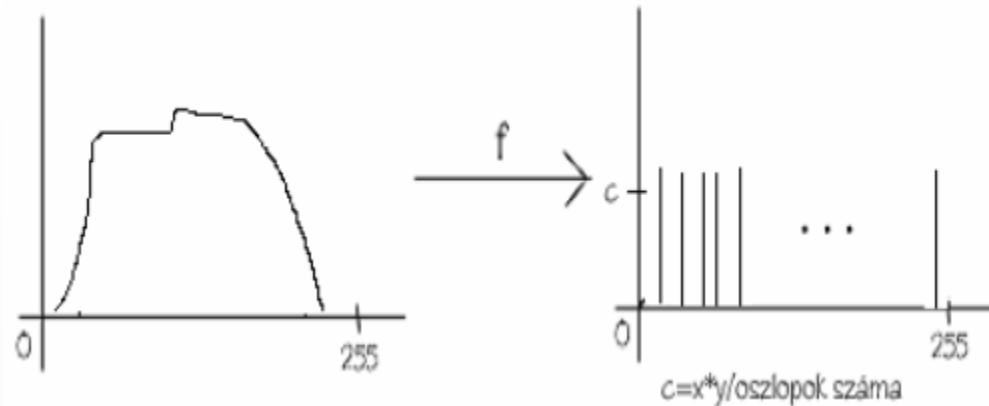
- Task: converting a grayscale image to binary (black&white).
What is the optimal threshold value?

- A possible good solution is to prescribe that the number of black and white pixels should be approximately the same in the output image.
 - The g_t threshold value can be calculated from the histogram, (P is the total number of pixels):

$$\sum_{i=0}^{g_t} h[i] \approx \sum_{i=g_t+1}^{255} h[i] \approx \frac{P}{2}$$

Generalization: histogram equalization

- **Goal:** contrast enhancement
- **Transform:** step (staircase) function. The number of columns determines number of color (intensity) values appearing in the output, (e.g. number of columns=16, 32, 64, etc.).



Histogram equalization – c output values

- Goal: determining the $t_0=0, t_1, \dots t_{c-1}, t_c=255$ dividing points, where:

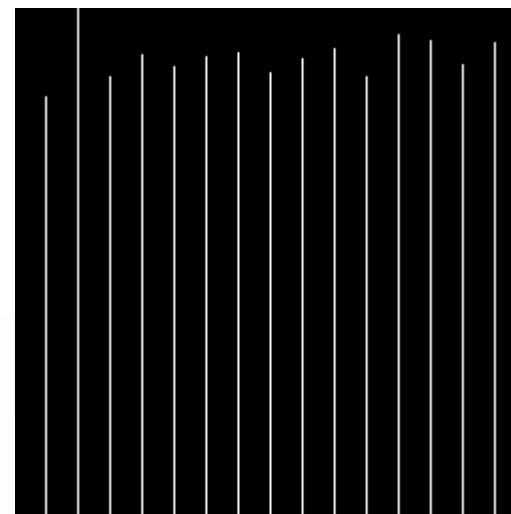
$$\sum_{i=0}^{t_j} h[i] \approx P \cdot \frac{j}{c} \quad j \in \{1\dots c\}$$

- c is the number of different gray levels in the output image ($c=2$ for thresholding, but it can also be 16, 32, ... 256 as well)
- P is the total number of pixels again.

Histogram equalization - result



16 level ouptut



Histogram of the output image

Histogram Transformations

○ Adaptive Histogram Equalization:

- applies histogram equalization on parts of the image (called tiles) independently
- Use post processing to reduce artifacts at the borders of the tiles.



[1] Zuiderveld, Karel. "Contrast Limited Adaptive Histogram Equalization." *Graphic Gems IV*. San Diego: Academic Press Professional, 1994. 474–485.

Spatial Filtering

◎ Smoothing:

- Reduce the noise that may corrupt the image.

◎ A few noise types we will work with:

- Impulse noise, (aka salt and pepper noise)
- Additive Gaussian Noise



Additive Gaussian Noise



Impulse Noise

The windmill at Wijk bij Duurstede by Jacob van Ruisdael (1670)

Spatial Filtering

◎ Gaussian Smoothing:

- With $\sigma=0.75$



Original image

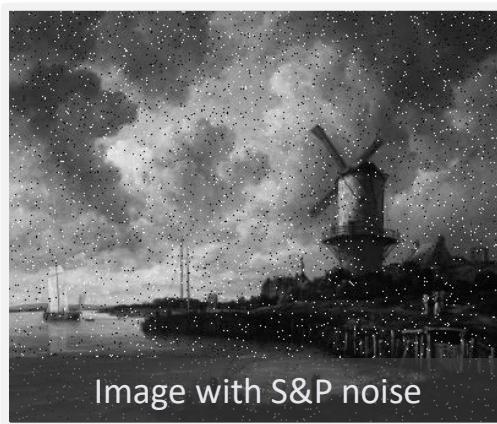


Image with S&P noise

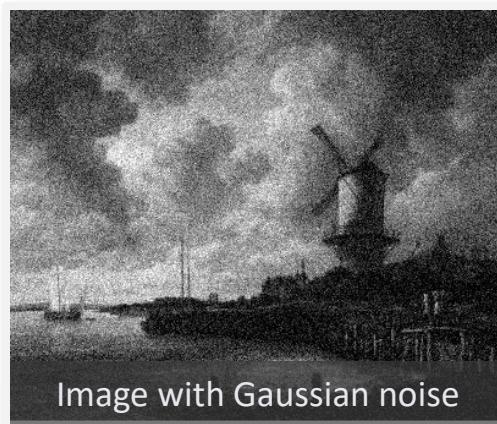
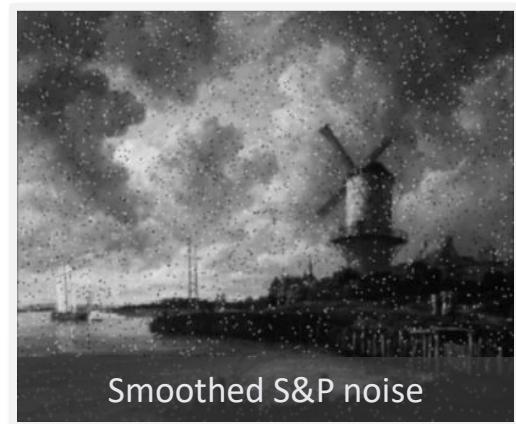


Image with Gaussian noise



Smoothed S&P noise



Smoothed Gaussian noise

Spatial Filtering

◎ Gaussian Smoothing:

- With $\sigma=1.5$



Original image

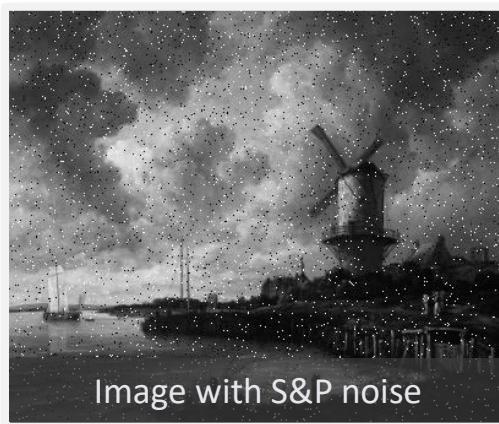


Image with S&P noise



Smoothed S&P noise

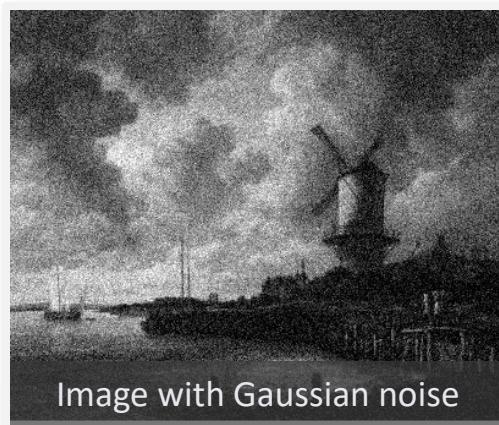


Image with Gaussian noise



Smoothed Gaussian noise

Spatial Filtering

◎ Spatially Adaptive Noise Smoothing:

- The smoothing takes into account the local characteristics of the image:

$$y(n_1, n_2) = \left(1 - \frac{\sigma_n^2}{\sigma_l^2}\right) \cdot x(n_1, n_2) + \frac{\sigma_n^2}{\sigma_l^2} \bar{x}(n_1, n_2)$$

$$\sigma_l^2(n_1, n_2) = \sum_{(n_1, n_2) \in N} (x(n_1, n_2) - \bar{x}(n_1, n_2))^2$$

Local variance of the image

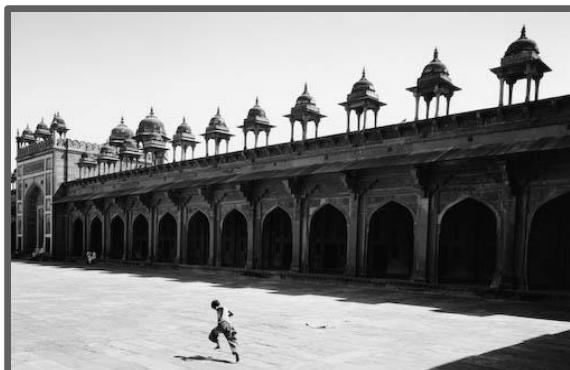
$$\bar{x}(n_1, n_2) = \frac{1}{|N|} \sum_{(n_1, n_2) \in N} x(n_1, n_2)$$

Local average of the image

Variance of the noise: either known a priori, or has to be measured

Spatial Filtering

◎ Spatially Adaptive Noise Smoothing:



Original image*

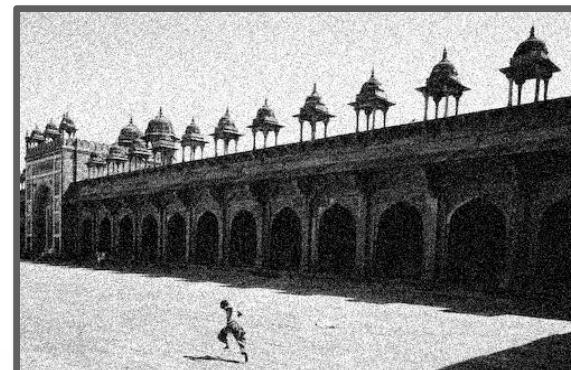
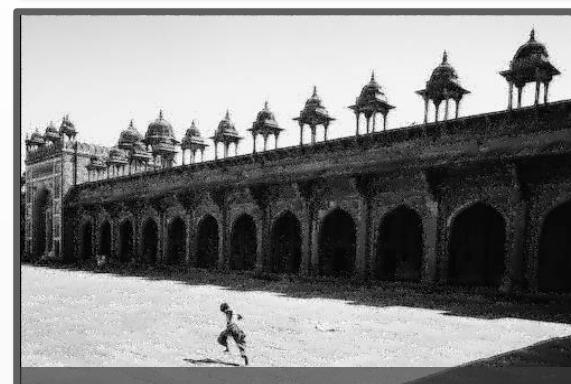


Image with Gaussian Noise



Gaussian Smoothing



Spatially Adaptive Smoothing

* Fatehpur Sikri, Inde by Lucien Hervé (1955)

Spatial Filtering

- **Rank filter:**

1. Consider the actual pixel and its neighborhood (e.g. $3 \times 3 = 9$ pixel sized window),
2. Sort the observed pixel values according to gray level,
3. Take the k -th value from this row as the new pixel value

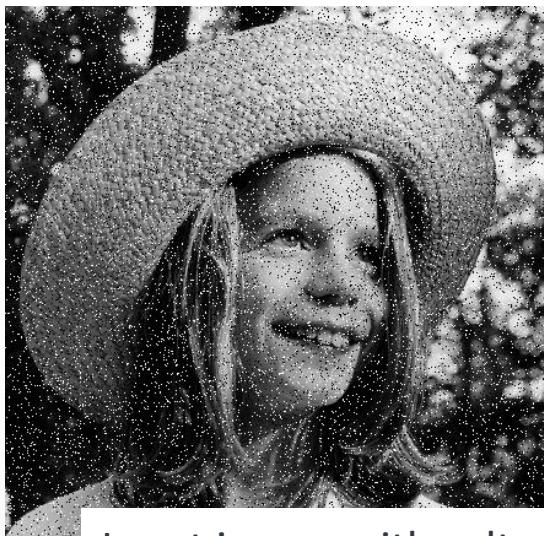
- **Median filter:** k is the middle pixel value in the row:

- $k = [(2W+1)^2 - 1]/2$, if W is the half side size of the neighborhood

- **Non-Linear filter**

Spatial Filtering

- **Median filter:** replaces each pixel with the *median value* of its analyzed neighborhood. (Median value: the center element of sorted values)
 - Very effective against impulse („salt and pepper“) noise:



Input image with salt
and pepper noise



Blur with convolution



Median filter

Spatial Filtering

- ◎ **Median filter:** replaces each pixel with the *median value* of its analyzed neighborhood. (Median value: the center element of sorted values)
 - Very effective against impulse („salt and pepper“) noise:



Original image



Image with S&P noise



Median filtered S&P noise

- Not so effective against Gaussian noise.

Spatial Filtering

○ Order statistic filtering:

- Based on the sorted pixel intensity levels in the analyzed neighborhood.
- If after sorting...
 - we take the middle element, we get back the median filter.
 - We take the maximum element to filter „pepper” and min to filter „salt” noise.



- But max filter will highlight „salt”, while min filter will highlight „pepper”.

Spatial Filtering

◎ Order statistic filtering:

- Mid-point filtering:
 - works well on Gaussian or uniform noise

$$y(n_1, n_2) = \frac{1}{2} \left(\max_{(m_1, m_2) \in N} \{x(m_1, m_2)\} + \min_{(m_1, m_2) \in N} \{x(m_1, m_2)\} \right)$$

- Alpha-trimmed mean filter:

$$y(n_1, n_2) = \frac{1}{|N| - \alpha} \sum_{(m_1, m_2) \in N_r} x(m_1, m_2)$$

- Where N_r is a reduced neighborhood, not containing the lowest and highest α element of N .
- If $\alpha = 0$, we get back the arithmetic mean.
- If $\alpha = |N| - 1$, we get back the median filter.

Wallis Operator

- The Wallis operator can help to adjust local contrast:

$$y(n_1, n_2) = [x(n_1, n_2) - \bar{x}(n_1, n_2)] \frac{A_{\max} \sigma_d}{A_{\max} \sigma_l(n_1, n_2) + \sigma_d} + [p\bar{x}_d + (1-p)\bar{x}(n_1, n_2)]$$

- where σ_l the local contrast: $\sigma_l(n_1, n_2) = \frac{1}{|N|} \sqrt{\sum_{(n_1, n_2) \in N} (x(n_1, n_2) - \bar{x}(n_1, n_2))^2}$
- \bar{x} is the local average: $\bar{x}(n_1, n_2) = \frac{1}{|N|} \sum_{(n_1, n_2) \in N} x(n_1, n_2)$
- σ_d is the desired local contrast, \bar{x}_d is the desired mean value of all pixels, p is a weighting factor of the mean compensation, while A_{\max} is maximizing the local contrast modification.

Wallis Operator

- We can describe the image the following way:

$$x(n_1, n_2) = \underbrace{[x(n_1, n_2) - \bar{x}(n_1, n_2)]}_{(1)} + \underbrace{\bar{x}(n_1, n_2)}_{(2)}$$

where (2) is the local mean and (1) is the deviation from the local mean.

- With the transformation we want to „push” the local mean and standard deviation to a predefined desired value:

$$y(n_1, n_2) = \underbrace{[x(n_1, n_2) - \bar{x}(n_1, n_2)]}_{(1)} \frac{\sigma_d}{\sigma_l(n_1, n_2)} + \underbrace{[p\bar{x}_d + (1-p)\bar{x}(n_1, n_2)]}_{(2)}$$

- We are almost there, but if the local contrast is too low, the weighting in (1) may get too high, this is why we maximize it with A_{max} :

$$\frac{\sigma_d}{\sigma_l(n_1, n_2)} \Rightarrow \frac{A_{max}\sigma_d}{A_{max}\sigma_l(n_1, n_2) + \sigma_d}$$

Wallis Operator



Original Image



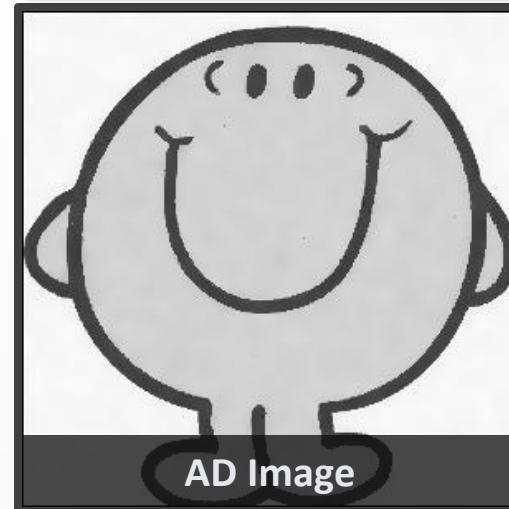
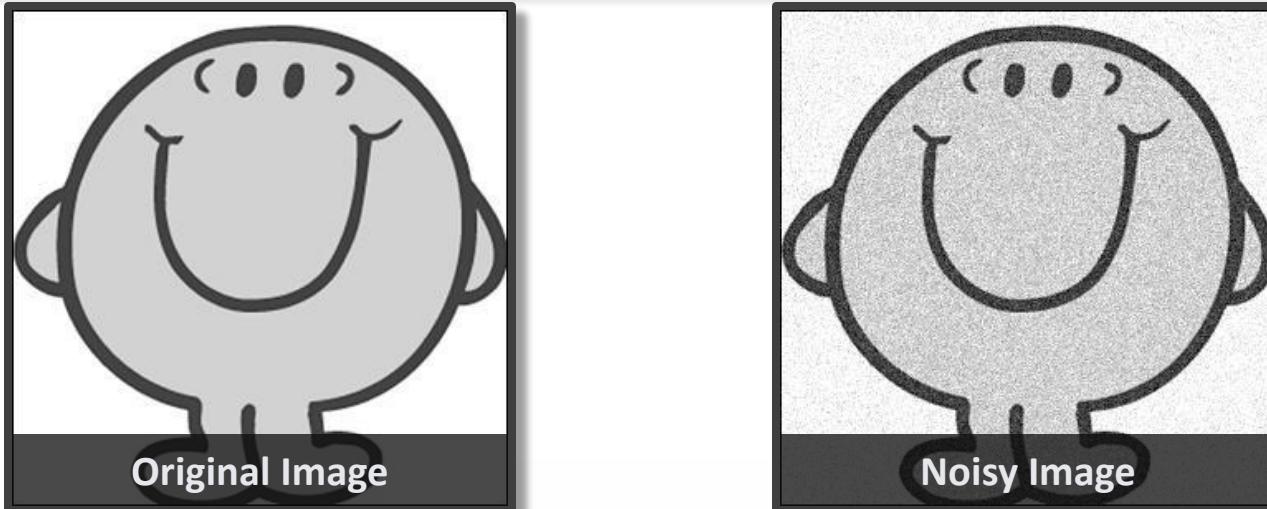
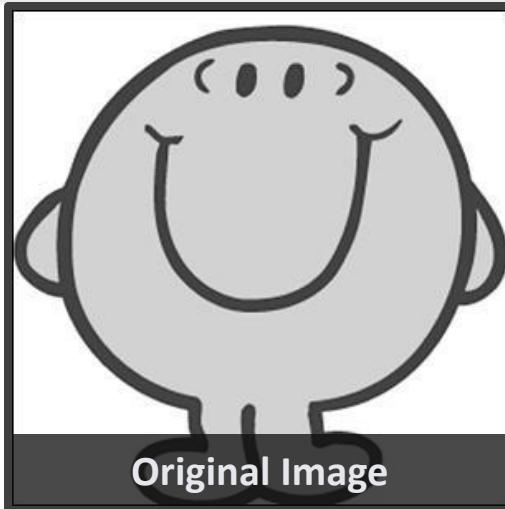
Image after applying Wallis operator

Anisotropic Diffusion

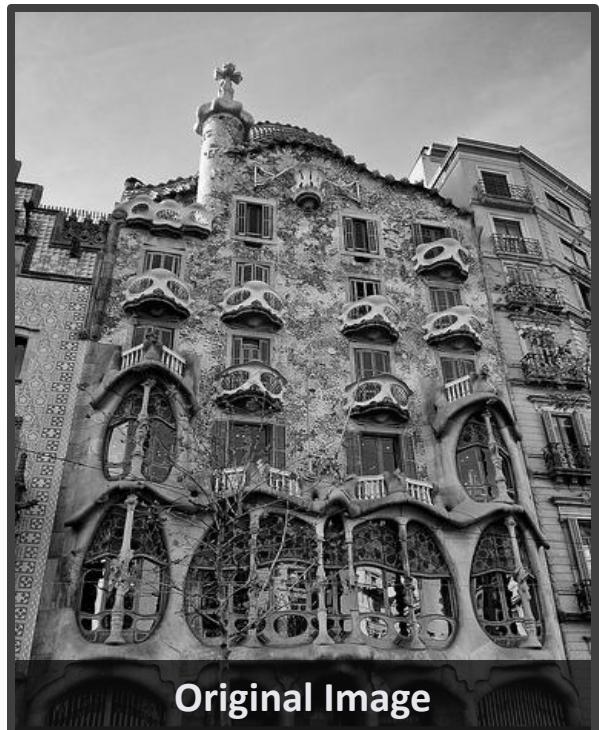
- The anisotropic diffusion is a technique aiming at reducing image noise without blurring significant parts of the image content.
- It was first proposed by D. Gábor in 1965 and later by Perona and Malik around 1990.
- ***Non-linear*** and ***space-variant*** transformation.
- The main idea is that the effect of blurring in each direction is inversely proportional to the gradient value in that direction:
 - allows diffusion along the edges or in edge-free territories, but penalizes diffusion orthogonal to the edge direction.
- AD is an iterative process

P. Perona, J Malik (July 1990). "Scale-space and edge detection using anisotropic diffusion". IEEE Tr. PAMI, 12 (7): 629–639.
D. Gabor, "Information theory in electron microscopy," Laboratory Investigation, vol. 14/6, pp. 801–807, 1965.

Anisotropic Diffusion



Anisotropic Diffusion



Original Image



Noisy Image



AD Image

Total Variation Regularization

- Assumption:

- The image is smooth inside the objects, with jumps across the boundaries.
- The noise component has high variation.

- The goal of ***Total Variation based noise removal*** is to minimize the total variation of the image while keep the result as close to the original input image as possible.

- It was introduced by Rudin, Osher and Fatemi in 1992.

Rudin, L. I.; Osher, S.; Fatemi, E. (1992). "Nonlinear total variation based noise removal algorithms". *Physica D* **60**: 259–268

Total Variation Regularization

- TV of the output image y is defined as the integral of the absolute gradient of the signal:

$$V(y) = \sum_{n_1} \sum_{n_2} \sqrt{|y(n_1+1, n_2) - y(n_1, n_2)|^2 + |y(n_1, n_2+1) - y(n_1, n_2)|^2}$$

- On the other hand, we also measure the difference between the original image x and the output image y by L_2 norm E :

$$E(x, y) = \sum_{n_1, n_2} (x(n_1, n_2) - y(n_1, n_2))^2$$

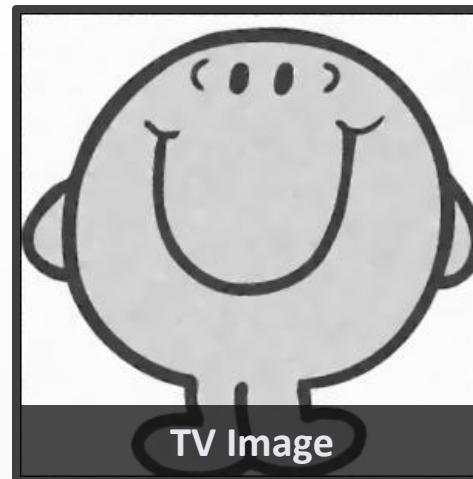
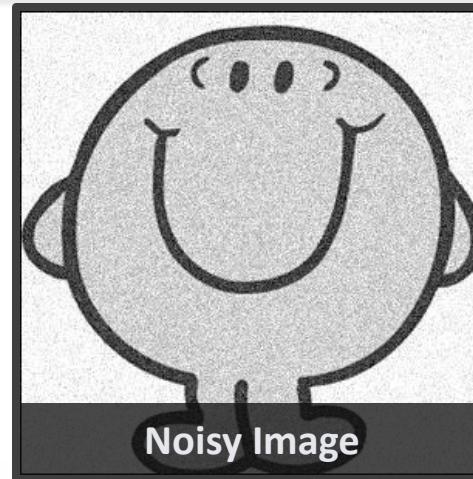
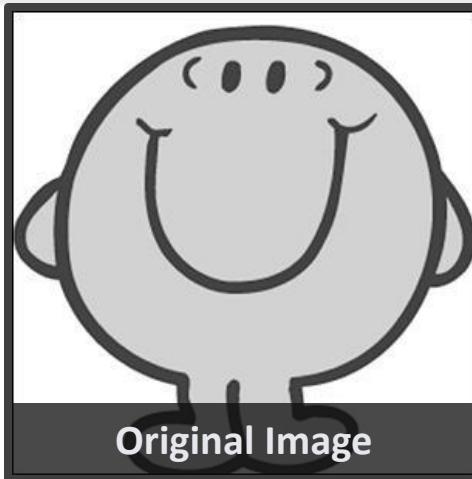
- The goal function for Total Variation based regularization:

$$\hat{y} = \arg \min_y [E(x, y) + \lambda V(y)]$$

where λ is the regularization parameter.

Rudin, L. I.; Osher, S.; Fatemi, E. (1992). "Nonlinear total variation based noise removal algorithms". *Physica D* **60**: 259–268

Total Variation Regularization

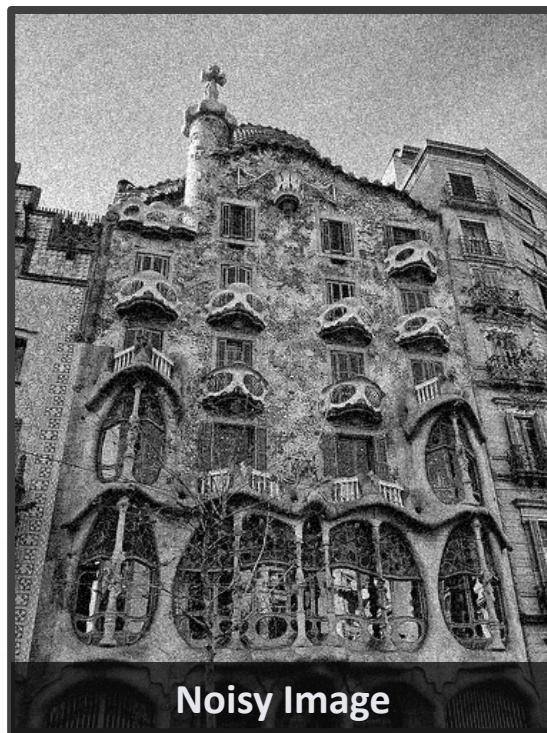


Matlab Code: http://visl.technion.ac.il/~gilboa/PDE-filt/tv_denoising.html

Total Variation Regularization



Original Image



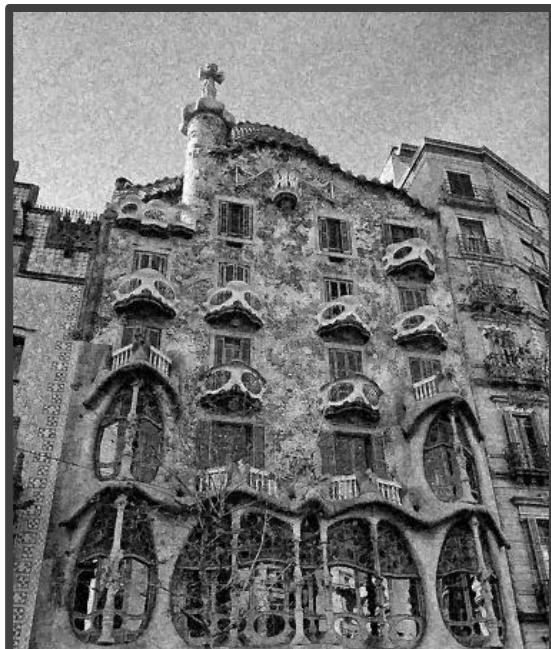
Noisy Image



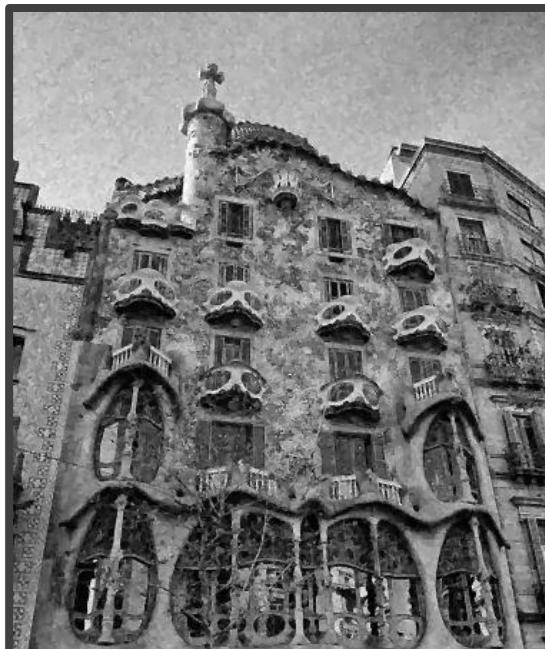
TV Image

Matlab Code: http://visl.technion.ac.il/~gilboa/PDE-filt/tv_denoising.html

Total Variation Regularization



TV Image (#iter=20)



TV Image (#iter=50)



TV Image (#iter=100)

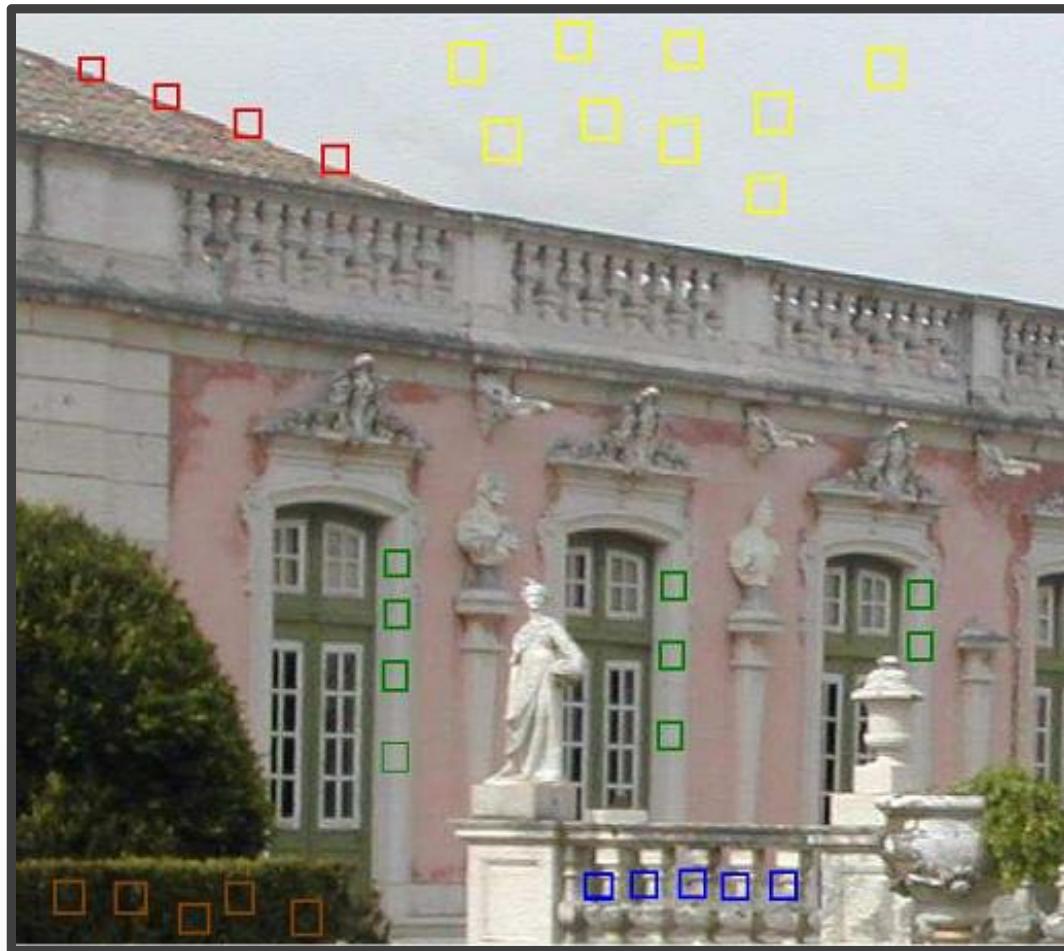
Matlab Code: http://visl.technion.ac.il/~gilboa/PDE-filt/tv_denoising.html

Non-Local Means Denoising

- The local smoothing methods aim at a noise reduction and at a reconstruction of the main geometrical configurations but not at the preservation of the fine structure, details and texture.
- ***High frequency image components are removed along with the noise, because they behave in all functional aspects as noise.***
- The NL-means algorithm tries to take advantage of the ***high degree of redundancy of any natural image:***
 - Every small window in a natural image has many similar windows in the same image.
 - The NL-means algorithm estimates the value of a pixel x as an average of the values of all the pixels whose Gaussian neighborhood looks like the neighborhood of x .

More details and a demo is available here: http://www.ipol.im/pub/art/2011/bcm_nlm/

Non-Local Means Denoising

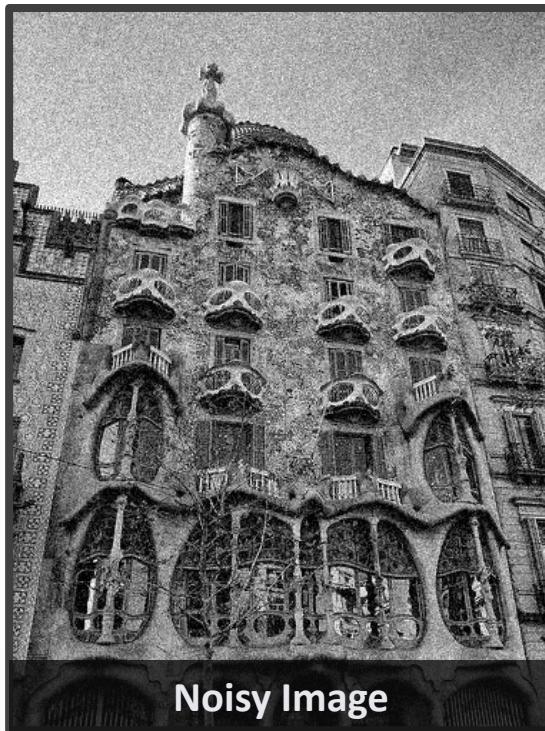


A. Efros and Th. K. Leung: Texture Synthesis by Non-parametric Sampling, IEEE ICCV, Corfu, Greece, September 1999

Non-Local Means Denoising



Original Image



Noisy Image



NLM Image

More details and a demo is available here: http://www.ipol.im/pub/art/2011/bcm_nlm/

Non-Local Means Denoising



Original Image



Noisy Image



NLM Image

More details and a demo is available here: http://www.ipol.im/pub/art/2011/bcm_nlm/