

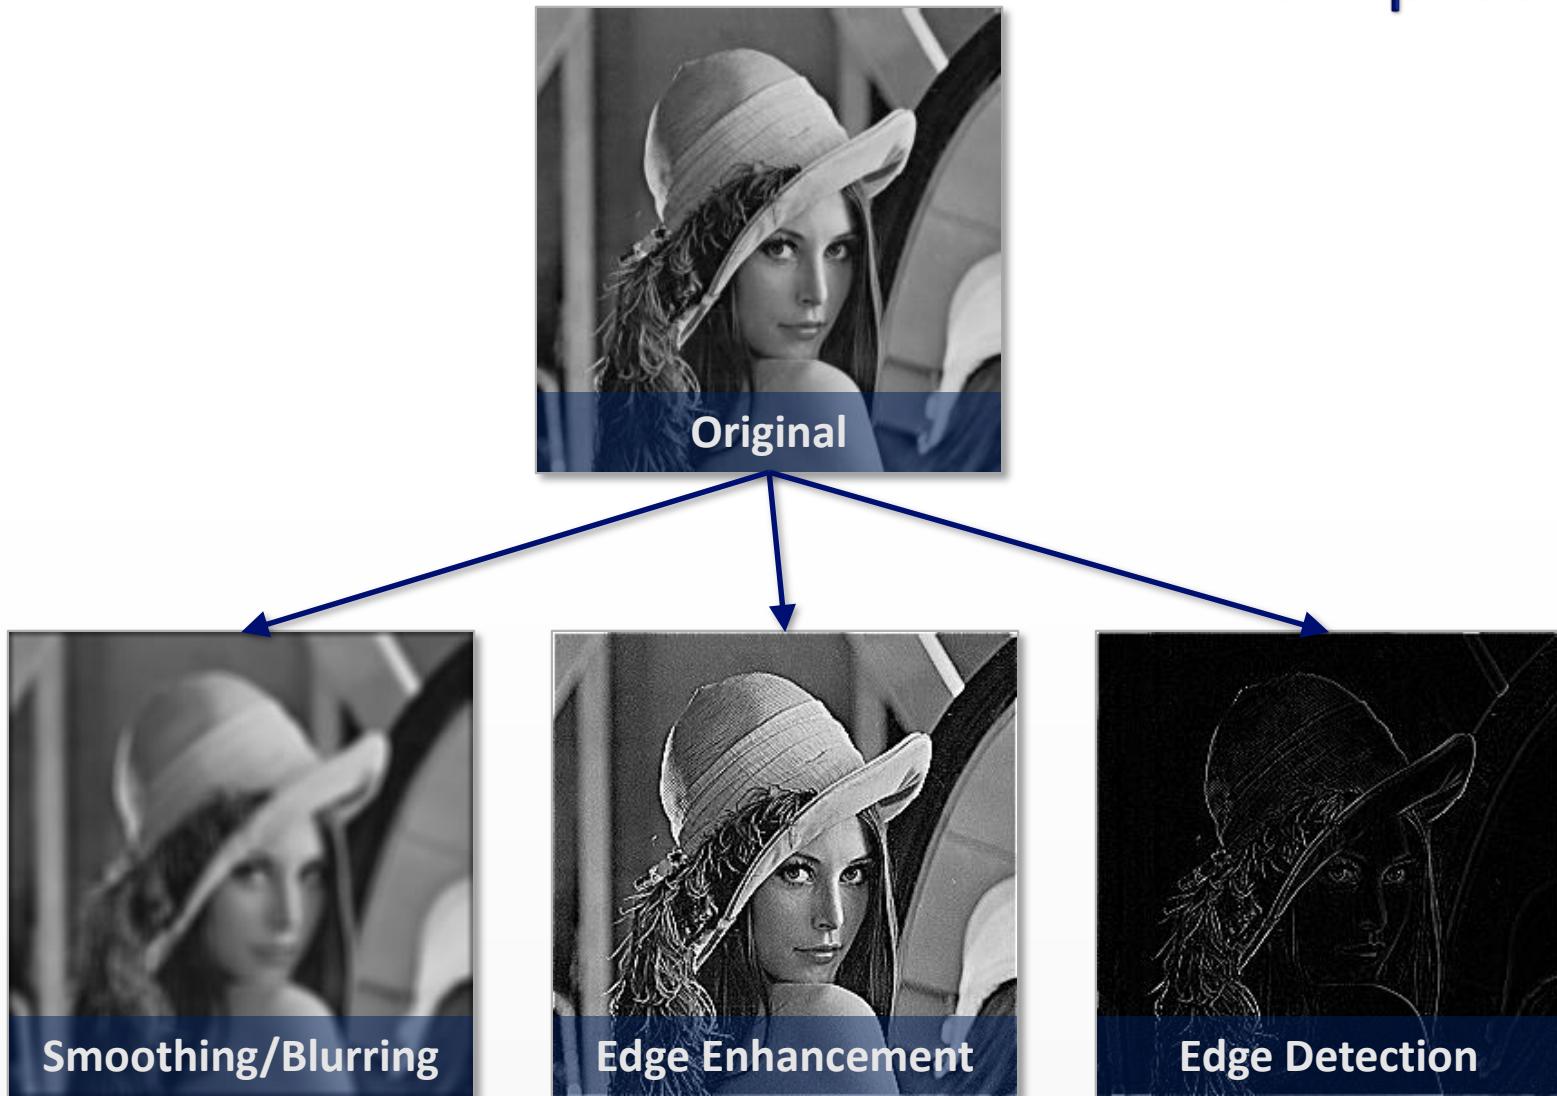
Basic Image Processing Algorithms

PPKE-ITK

Lecture 2.

2D Convolution

Examples



Mathematical background

- We look at the image as a 2D function:
- We can define different transformations:

- Intensity value inversion:

$$y(n_1, n_2) = 255 - x(n_1, n_2)$$

- Intensity shift with constant:

$$y(n_1, n_2) = x(n_1, n_2) + 100$$

- Weighting:

$$y(n_1, n_2) = x(n_1, n_2)w(n_1, n_2)$$

- Average on an N neighborhood:

$$y(n_1, n_2) = \text{average } \left(N(x(n_1, n_2)) \right)$$

- 2 important properties of the transformations we want to use on images: ***linearity*** and ***shift invariance***

Mathematical background

◎ Linearity:

$$T[x_1(n_1, n_2) + x_2(n_1, n_2)] = T[x_1(n_1, n_2)] + T[x_2(n_1, n_2)]$$

$$T[\alpha x(n_1, n_2)] = \alpha T[x(n_1, n_2)]$$

- e.g.: weighting is linear, intensity inversion is non-linear

◎ Spatial Invariance:

$$T[x(n_1, n_2)] = y(n_1, n_2)$$

$$T[x(n_1 - k_1, n_2 - k_2)] = y(n_1 - k_1, n_2 - k_2)$$

- e.g.: weighting is not SI, intensity inversion is SI
- e.g.: averaging on neighborhood is LSI

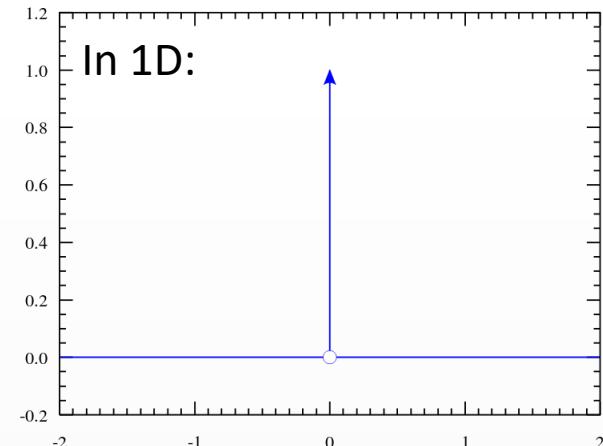
Unit Impulse Function

- ◎ 2D Unit Impulse function (Delta function) on \mathbb{Z} as follows:

$$\delta(n_1, n_2) = \begin{cases} 1 & \text{when } n_1 = 0 \text{ and } n_2 = 0 \\ 0 & \text{otherwise} \end{cases}$$

- ◎ For any 2D function $x(n_1, n_2)$:

$$x(n_1, n_2) = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} \delta(n_1 - k_1, n_2 - k_2) x(k_1, k_2)$$



Convolution

- ◎ **Impulse response** is the output of an LSI transformation if the input was the Delta function: $\delta(n_1, n_2) \rightarrow T \rightarrow h(n_1, n_2)$

If T is an LSI system:

$$T[x(n_1, n_2)] = y(n_1, n_2)$$

Then we can define convolution as follows:

$$\begin{aligned}y(n_1, n_2) &= x(n_1, n_2) * h(n_1, n_2) = \\&= h(n_1, n_2) * x(n_1, n_2) = \\&= \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x(k_1, k_2)h(n_1 - k_1, n_2 - k_2)\end{aligned}$$

Derivation of Convolution

$$y(n_1, n_2) = T[x(n_1, n_2)] =$$

$$= T \left[\sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x(k_1, k_2) \delta(n_1 - k_1, n_2 - k_2) \right] =$$

Linearity

$$= \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x(k_1, k_2) T[\delta(n_1 - k_1, n_2 - k_2)] =$$

Spatial Invariance

$$= \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x(k_1, k_2) h(n_1 - k_1, n_2 - k_2)$$

The Properties of Convolution

- ◎ Commutative:

$$f * g = g * f$$

- ◎ Associative:

$$f * (g * h) = (f * g) * h$$

- ◎ Distributive:

$$f * (g + h) = f * g + f * h$$

- ◎ Associative with scalar multiplication:

$$\alpha(f * g) = (\alpha f) * g$$

2D convolution in Practice

- In practice both the kernel and the image have finite size.

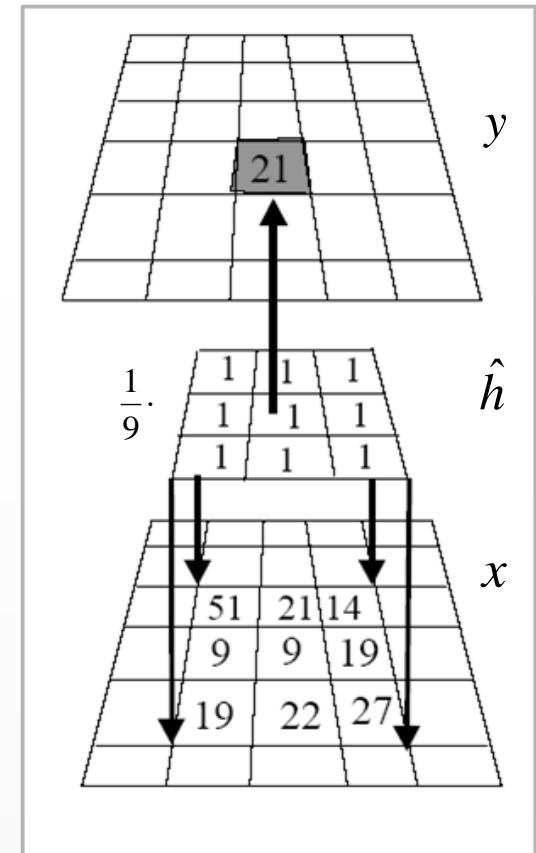
Let h and \hat{h} be $(2r_1 + 1) \times (2r_2 + 1)$ sized kernels, where \hat{h} is the 180° rotated version of h :

$$h = \begin{bmatrix} a_{-r_1, -r_2} & \cdots & a_{-r_1, r_2} \\ \vdots & \ddots & \vdots \\ a_{r_1, -r_2} & \cdots & a_{r_1, r_2} \end{bmatrix} \text{ and } \hat{h} = \begin{bmatrix} a_{r_1, r_2} & \cdots & a_{r_1, -r_2} \\ \vdots & \ddots & \vdots \\ a_{-r_1, r_2} & \cdots & a_{-r_1, -r_2} \end{bmatrix}$$

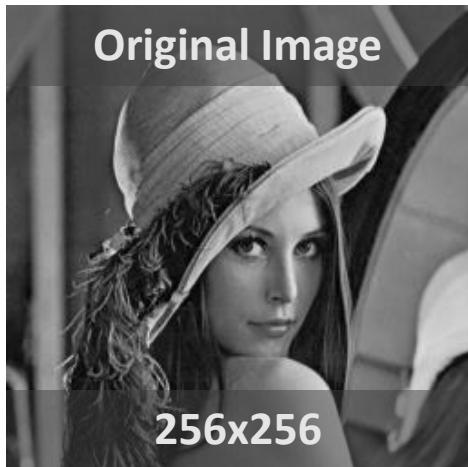
$$y(n_1, n_2) = \sum_{k_1=-r_1}^{r_1} \sum_{k_2=-r_2}^{r_2} x(k_1, k_2) h(n_1 - k_1, n_2 - k_2) =$$

$$= \sum_{k_1=-r_1}^{r_1} \sum_{k_2=-r_2}^{r_2} h(k_1, k_2) x(n_1 - k_1, n_2 - k_2) =$$

$$= \sum_{k_1=-r_1}^{r_1} \sum_{k_2=-r_2}^{r_2} \hat{h}(k_1, k_2) x(n_1 + k_1, n_2 + k_2)$$



Size of the Convolved Image



Convolutional Kernel

$$* \frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix} =$$

5x5



In general:

Size of the input image: $(A \times B)$

Size of the kernel: $(C \times D)$

Size of the output image: $(A+C-1) \times (B+D-1)$

Boundary Effects

- What happens at the border of the image?



Original image with
the problematic area



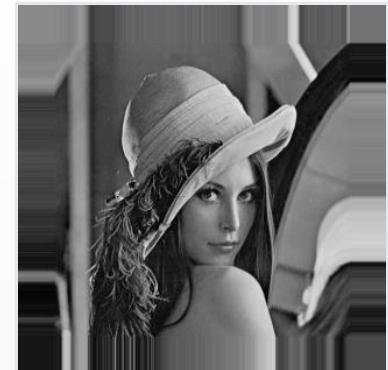
Zero padding



Mirroring



Circular padding



Repeating border

Applications

- Possible application of convolution:

- Smoothing/Noise reduction
- Edge detection
- Edge enhancement

- Depending on the task the *sum of the elements of the kernel matrix* can be different:

- 1: smoothing, edge enhancement

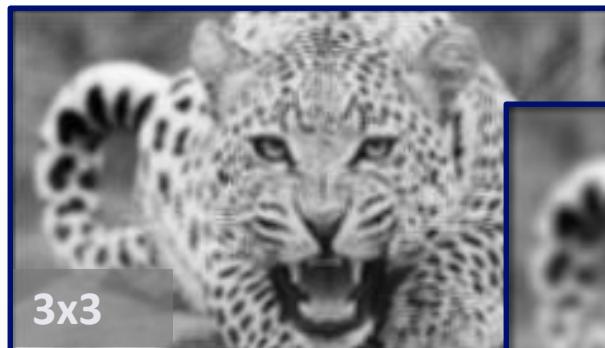
E.g.: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$

- 0: edge detection

E.g.: $\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$ $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$

Smoothing/Blurring

- Simple average:



Computational requirements

- For k_s kernel size and P image size (area, measured in pixels) approximately $\sim k_s P$ operations are needed.
- For large kernel size the execution may be slow

Decreasing the computational need for a simple (averaging) blur operation

- Integral image: $f \rightarrow I_f$
auxilliary representation

$$I_f(x, y) = \sum_{i=1}^x \sum_{j=1}^y f(i, j)$$

- E.g.: $I_f(3,3) = \text{sum of the values of pixels: } 11$

1	0	2	1
2	0	1	0
3	1	1	0
1	0	1	4



I_f

1	0	2	1
2	0	1	0
3	1	1	0
1	0	1	4

1	1	3	4
3	3	6	7
6	7	11	12
7	8	13	18

Calculation of I_f with dynamic programming in $\sim P$ time:

- Auxiliary-auxiliary image: $t(x, y) = \sum_{j=1}^y f(x, j)$

$$t(x, 1) := f(x, 1), x = 1 \dots w \quad t(x, y) = t(x, y - 1) + f(x, y)$$

$$I_f(1, y) := t(1, y), \quad y = 1 \dots h$$

$$I_f(x, y) = I_f(x - 1, y) + t(x, y)$$

$$\textcolor{green}{4} + \textcolor{red}{7} = \textcolor{blue}{11}$$

1	0	2	1
2	0	1	0
3	1	1	0
1	0	1	4

$t :$

1	0	2	1
3	0	3	1
6	1	4	1
7	1	5	5

I_f

1	1	3	4
3	3	6	7
6	7	11	12
7	8	13	18

$f :$

Utilization of the integral image

- Sum of pixel values in an **arbitrary sized** sub-rectangle can be calculated by applying **3 additive operations** using the integral image:

$$\sum_{i=a}^c \sum_{j=b}^d f(i, j) = I_f(c, d) - I_f(a-1, d) - I_f(c, b-1) + I_f(a-1, b-1)$$

- Example ($a=1, b=1, c=2, d=2$): $11-6-3+1=3$

	1	0	2	1
f	2	0	1	0
	3	1	1	0
	1	0	1	4

I_f

1	1	3	4
3	3	6	7
6	7	11	12
7	8	13	18

Using integral image for quick blurring (simple averaging kernel)

$$\tilde{f}(x, y) = \frac{1}{(2r+1)^2} \sum_{i=-r}^r \sum_{j=-r}^r f(x+i, y+j)$$

$(2r+1)^2$ addition
+ 1 division operations

Example: $r=5 \rightarrow$ For the whole image $\sim 122P$ operations

$$\begin{aligned} \tilde{f}(x, y) = & \frac{1}{(2r+1)^2} (I_f(x+r, y+r) - I_f(x-r-1, y+r) - \\ & - I_f(x+r, y-r-1) + I_f(x-r-1, x-r-1)) \end{aligned}$$

3 addition
+1 division

Example: $r=5 \rightarrow$ For the whole image $\sim 2P+4P=6P$ operations



Optional homework (a bit more than a convolution)

- Construct an efficient contrast calculating algorithm using the integral image! Contrast is calculated as the standard deviation of pixel values of the $(2r+1)^2$ size neighborhood of each pixel.

$$\sigma^2(x, y) = \frac{1}{(2r+1)^2} \sum_{i=-r}^r \sum_{j=-r}^r [f(x+i, y+j) - \tilde{f}(x, y)]^2$$



where: $\tilde{f}(x, y) = \frac{1}{(2r+1)^2} \sum_{i=-r}^r \sum_{j=-r}^r f(x+i, y+j)$

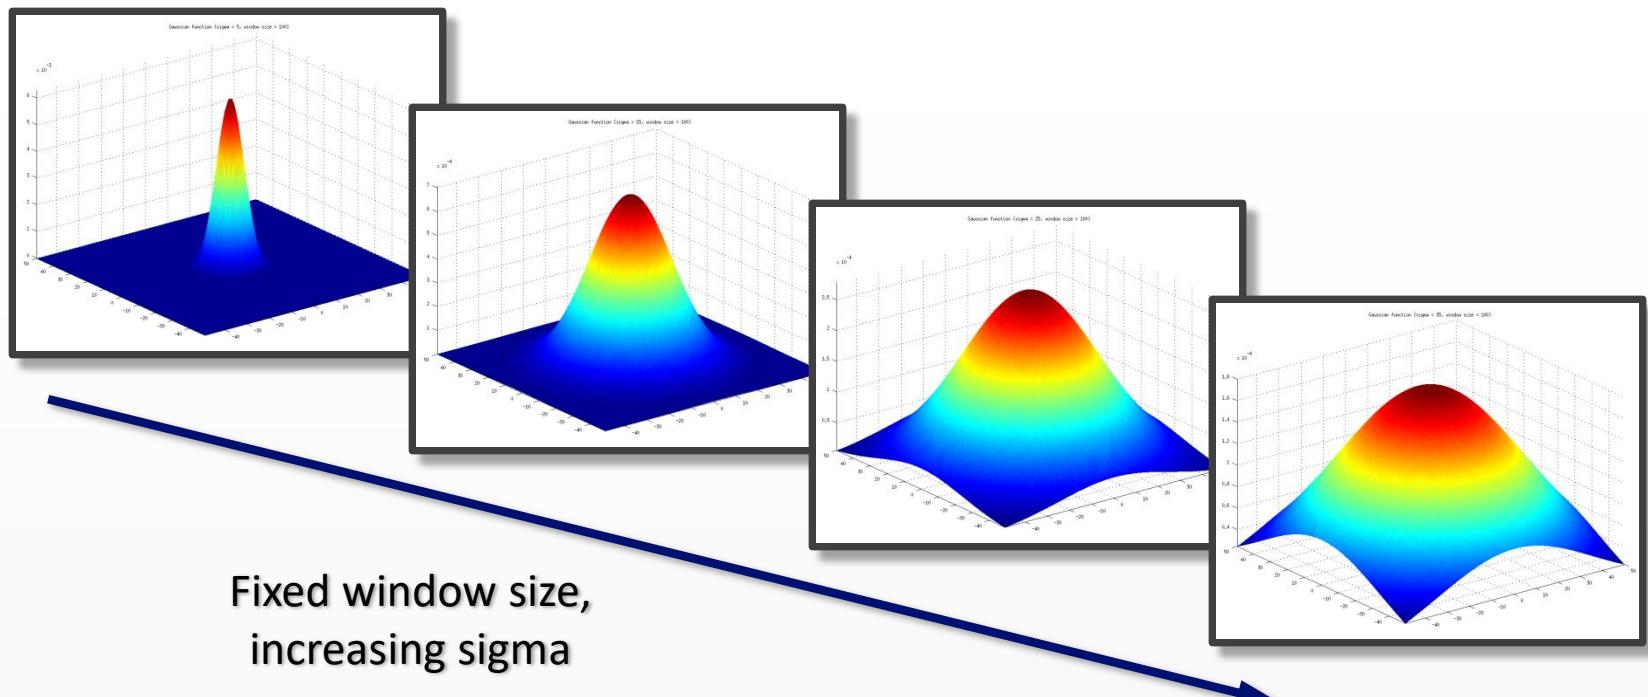
Hint:

$$\sigma^2(x, y) = \left\{ \frac{1}{(2r+1)^2} \sum_{i=-r}^r \sum_{j=-r}^r [f(x+i, y+j)]^2 \right\} - [\tilde{f}(x, y)]^2$$

Smoothing/Blurring

◎ Gaussian blur:

- Weights are defined by a 2D Gaussian function
- 2 parameters: **size of the window** and the **standard deviation** of the Gaussian



Smoothing/Blurring

◎ Gaussian blur:

- Weights are defined by a 2D Gaussian function
- 2 parameters: window size and the width of the Gaussian
- E.g. kernel size = 5x5; $\sigma = 1.5$;

$$\begin{bmatrix} 0.0144 & 0.0281 & 0.0351 & 0.0281 & 0.0144 \\ 0.0281 & 0.0547 & 0.0683 & 0.0547 & 0.0281 \\ 0.0351 & 0.0683 & 0.0853 & 0.0683 & 0.0351 \\ 0.0281 & 0.0547 & 0.0683 & 0.0547 & 0.0281 \\ 0.0144 & 0.0281 & 0.0351 & 0.0281 & 0.0144 \end{bmatrix}$$



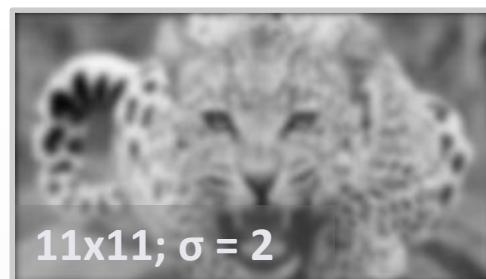
- E.g. kernel size = 3x3; $\sigma = 1.5$;

$$\begin{bmatrix} 0.0947 & 0.1183 & 0.0947 \\ 0.1183 & 0.1478 & 0.1183 \\ 0.0947 & 0.1183 & 0.0947 \end{bmatrix}$$

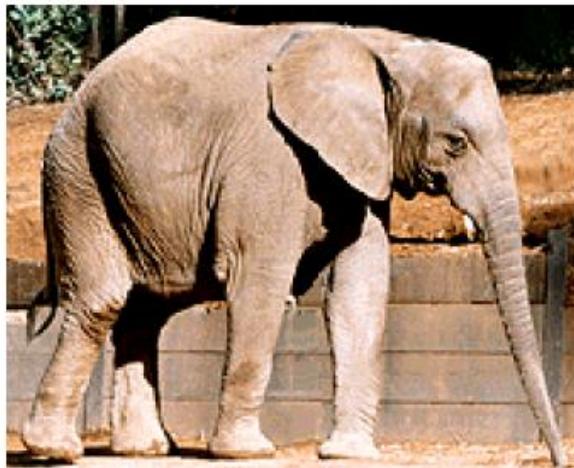
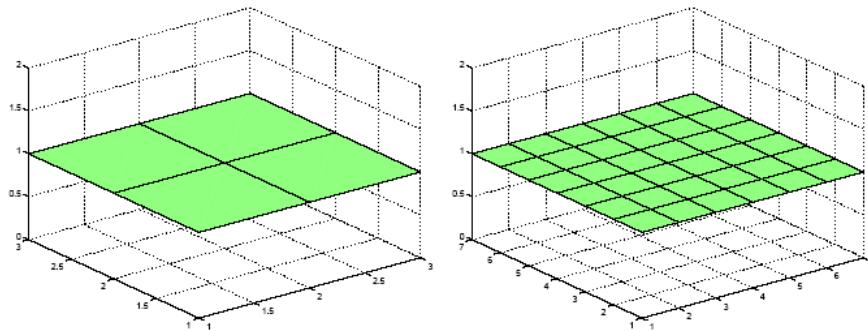


Smoothing/Blurring

- ◎ Gaussian blur:



Convolution examples – averaging blur

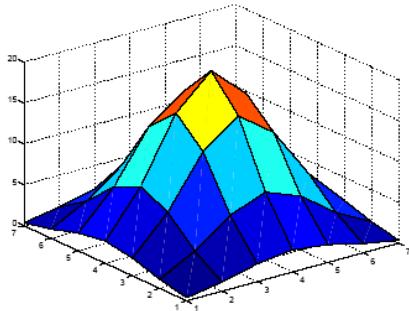
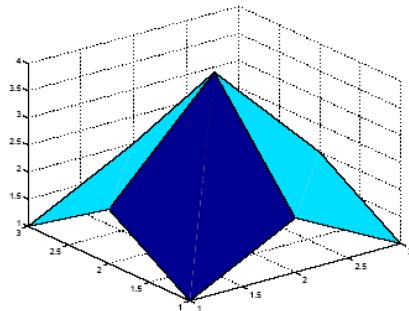


Input Image



Average blur

Convolution examples– Gaussian blur



$$K=1/123 \ast [1 \ 2 \ 3 \ 2 \ 1] \\ [2 \ 7 \ 11 \ 7 \ 2] \\ [3 \ 11 \ 17 \ 11 \ 3] \\ [2 \ 7 \ 11 \ 7 \ 2] \\ [1 \ 2 \ 3 \ 2 \ 1]$$



Input Image



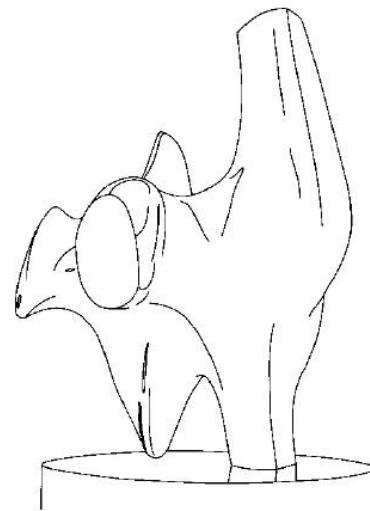
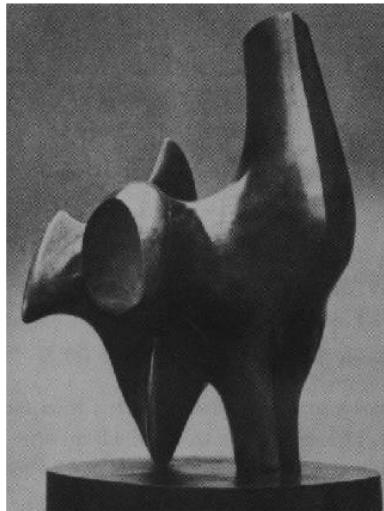
Gaussian blur

Edge detection

- Goal: extracting the object contours
- Edge points: brightness changes sharply



Goals of edge detection



- Goal: extracting curves from 2D images
 - More compact content representation than pixel
 - Segmentation, recognition, scratch filtering

Goals of edge detection

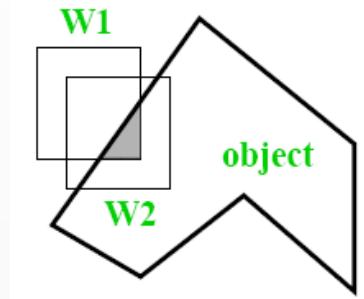
- Extracting image information, structures
 - Corners, lines, borders
- Not always simple...



Edge detection

- Properties of a good edge filter:

- (Near) zero output in homogeneous regions (constant intensity)
- Good detection :
 - detects as many real edges as possible
 - does not create false edges (because of e.g. image noise)
- Good localization: detected edges should be as close as possible to the real edges
- Isotropic: filter response independent on edge directions
 - all edges are detected regardless of their direction



Basic structures

- **Edge:** sharp intensity change (steep or continuous)
- **Line:** thin, long region with approx. uniform width and intensity level
- **Blob:** closed region with homogeneous intensity
- **Corner:** breaking or direction change of a contour or edge



Edge



Line



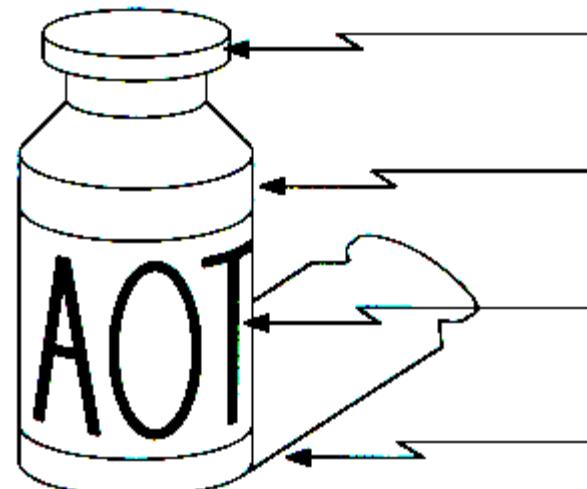
Blob



Corner

Origin of edges

- Various effects may cause edges



Sharp change in surface normals

Continuos change in surface depth

Change in surface color

Changes cased by illumination/shadows

Parameters of an edge

- Edge normal: vector, perpendicular to the edge, pointing toward the steepest intensity change
- Edge direction: vector point towards the direction of the line
 - Position, center point
- Strength: intensity ratio w.r.t. neighborhood

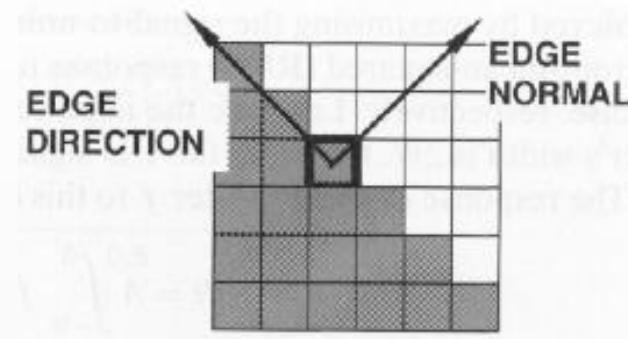
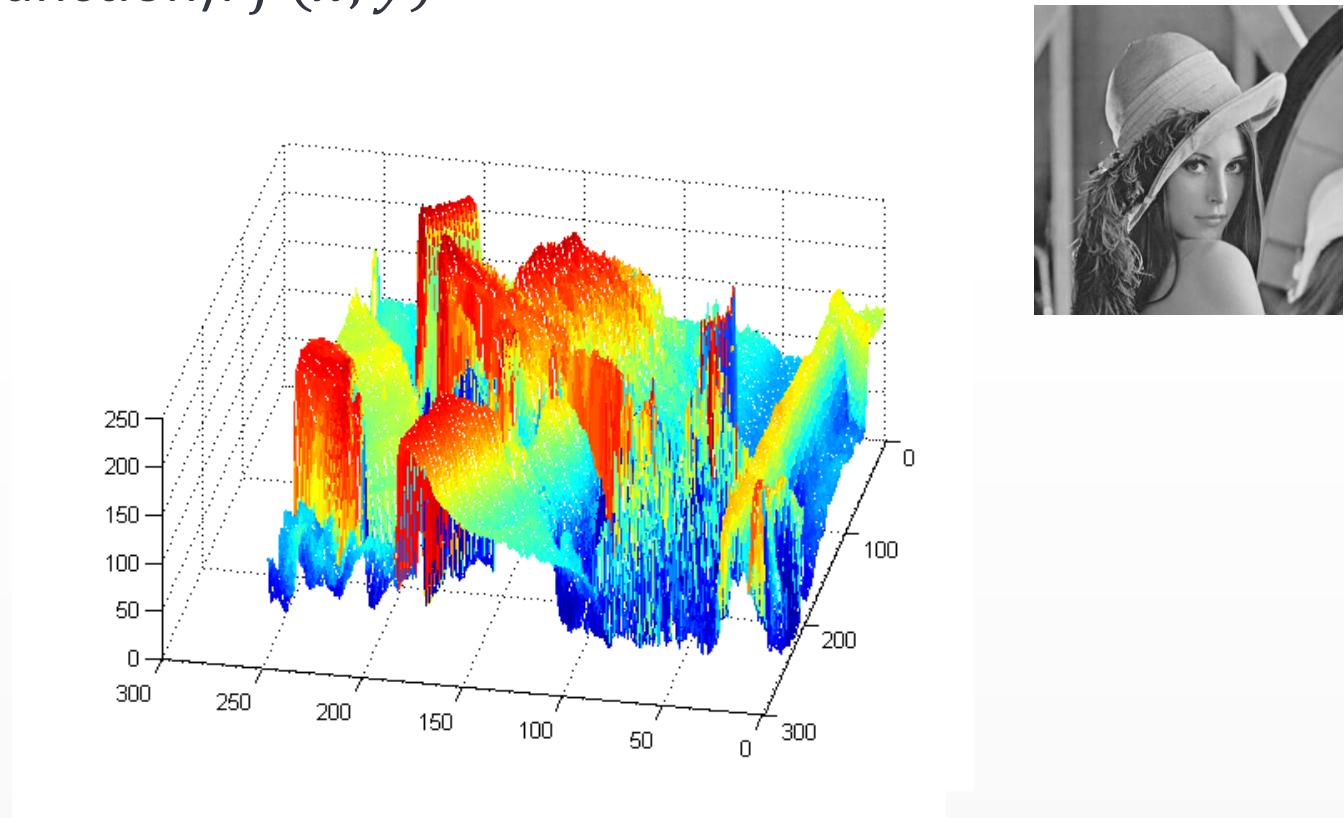


Image representation

- Image: gray value is function of the x and y coordinates
(intensity function): $f(x, y)$



Edge Detection

- **Edges types:** locations on the image where the *intensity changes sharply* (usually at the contour of objects)



- We are searching for places where the gradient of the 2D function (the image) is high.
- Main types of edge detection:
 - First order derivative
 - Second order derivative
 - Others:
 - Complex methods e.g. Canny method
 - Phase Congruency

Edge Detection

◎ Edge detection with first order derivative:

- Using the gradient vector:

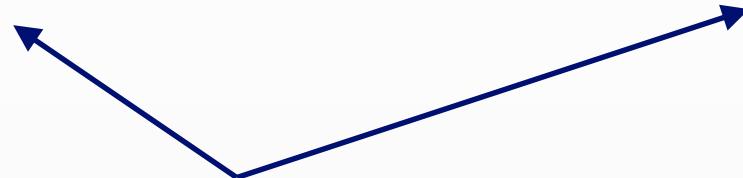
$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} \end{bmatrix}^T$$

- The approximation of the partial derivatives:

$$\frac{\partial f}{\partial x} = \lim \frac{f(x + dx, y) - f(x, y)}{dx} \quad \frac{\partial f}{\partial y} = \lim \frac{f(x, y + dy) - f(x, y)}{dy}$$

$$\approx f(x + 1, y) - f(x, y)$$

$$\approx f(x, y + 1) - f(x, y)$$



Since the smallest meaningful discrete value is $dx=1$ and $dy = 1$.

Edge Detection

◎ Edge detection with first order derivative:

- The approximation of the partial derivatives:

$$\frac{\partial f}{\partial x} \approx f(x+1, y) - f(x, y)$$

$$\frac{\partial f}{\partial y} \approx f(x, y+1) - f(x, y)$$

- Kernels for the gradient calculation with convolution (Prewitt):

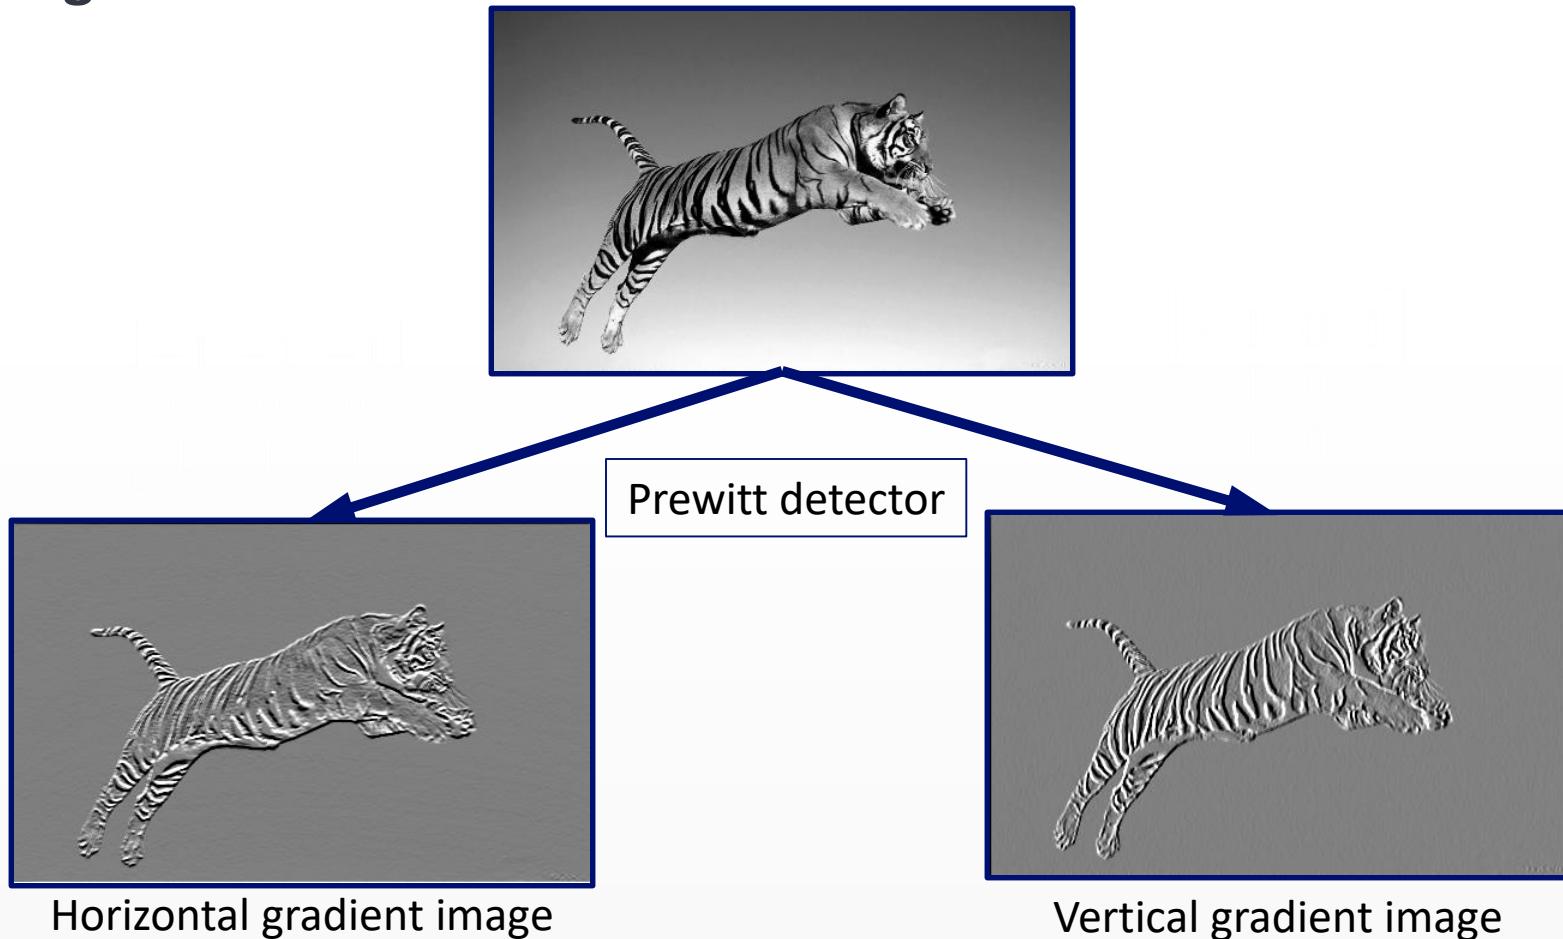
$$[-1 \ 1] \rightarrow [-1 \ 0 \ 1] \rightarrow [-1 \ 0 \ 1] * \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

For better localization For noise reduction

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix} \rightarrow \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix} \rightarrow \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix} * [1 \ 1 \ 1] = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

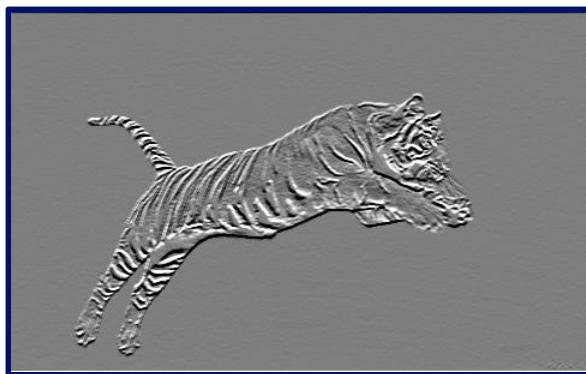
Edge Detection

- Edge detection with first order derivative:

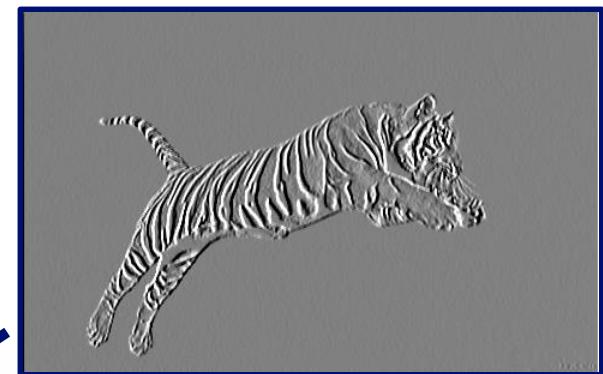


Edge Detection

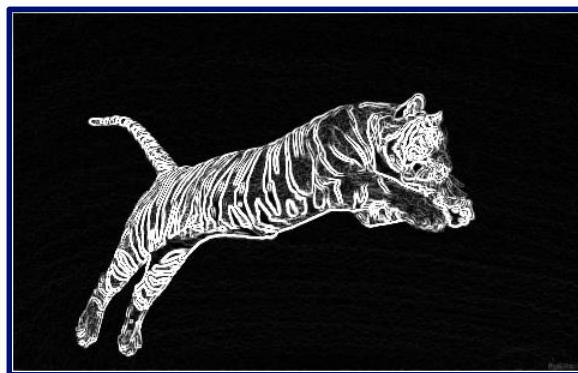
- Edge detection with first order derivative:



Horizontal gradient image



Vertical gradient image

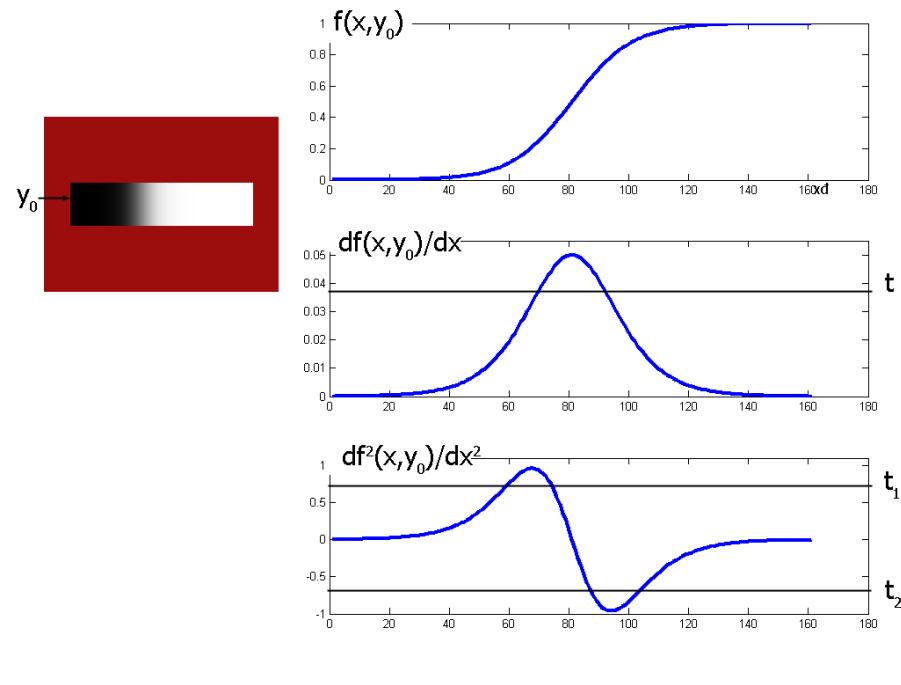


gradient image

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

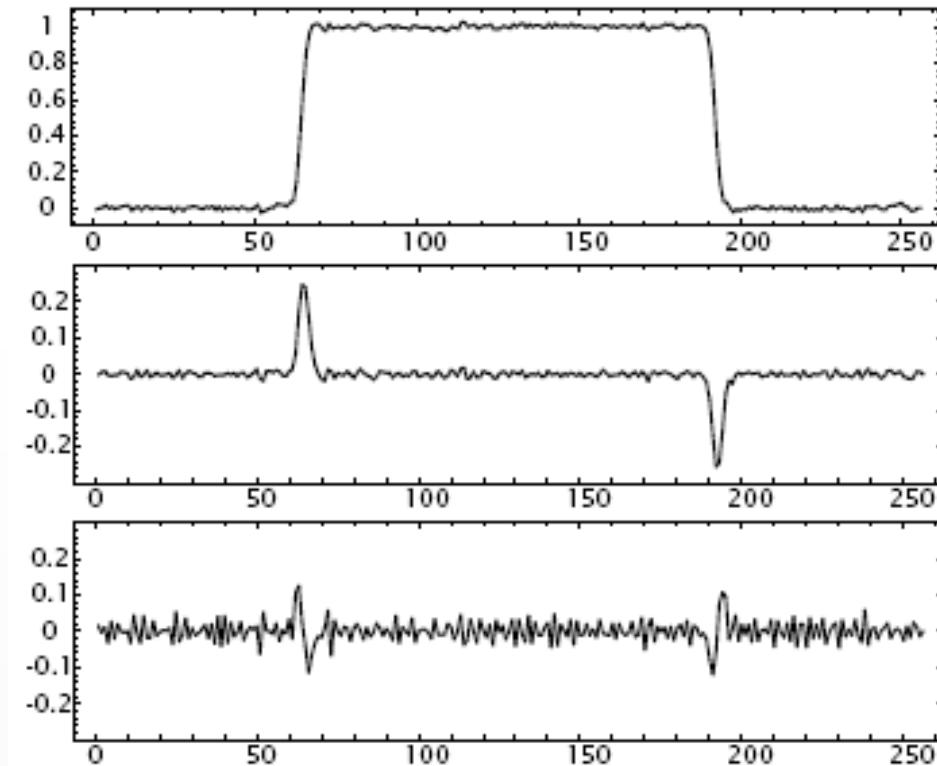
Second order edge detection: motivation

Horizontal edge detection
in the following image:



Top: intensity function along a selected horizontal line
Center: x directional first derivative
Bottom: x directional second derivative

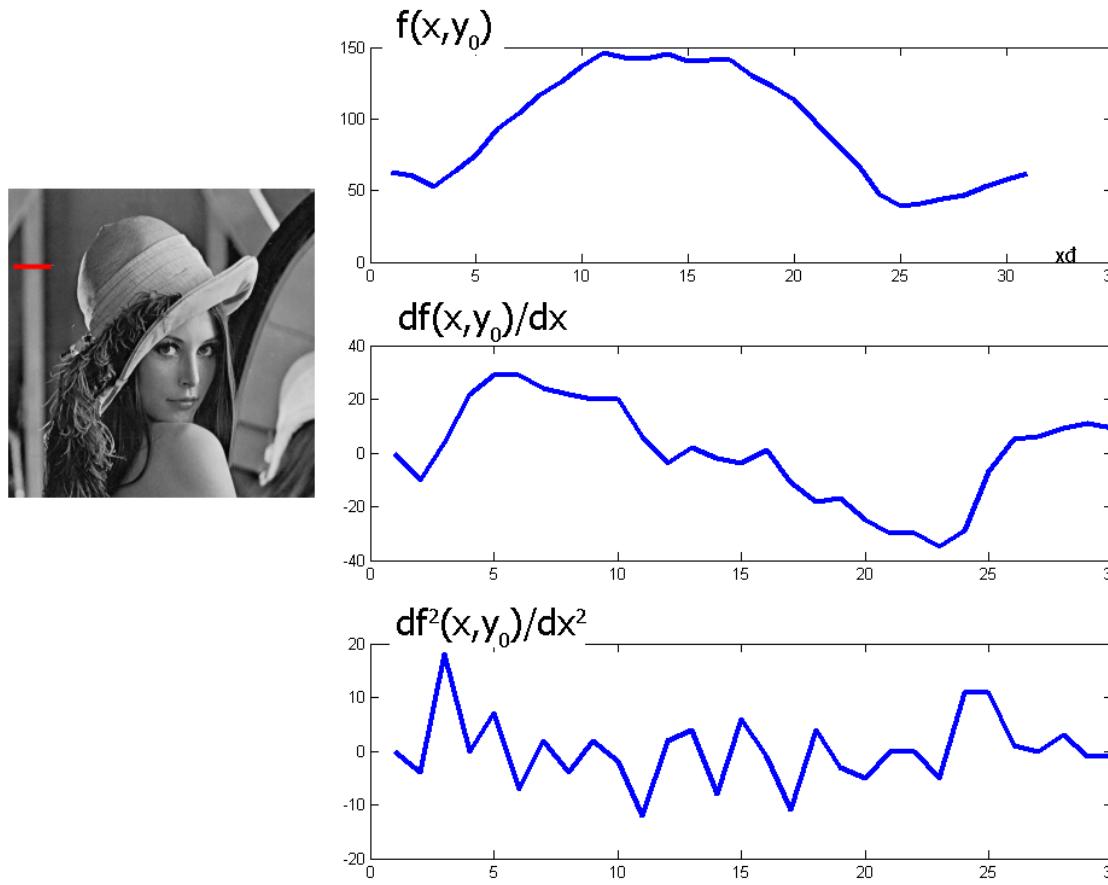
Second order case: instead of extreme values, search for zero crossing



Extreme values
(first order)

zero crossing
(second order)

Real photo: intensity profile below the red line segment



Edge Detection

◎ Edge detection with second order derivative:

- Approximation of the second order derivative for x direction:

$$\frac{\partial^2 f}{\partial x^2} = \lim_{d \rightarrow 0} \frac{\frac{\partial f}{\partial x}(x+d) - \frac{\partial f}{\partial x}(x)}{d} \approx \frac{\partial f}{\partial x}(x+1) - \frac{\partial f}{\partial x}(x) =$$

$$= f(x+2) - 2f(x+1) + f(x) \quad (\text{y direction similarly})$$

- Kernel for the second order gradient calculation with convolution:

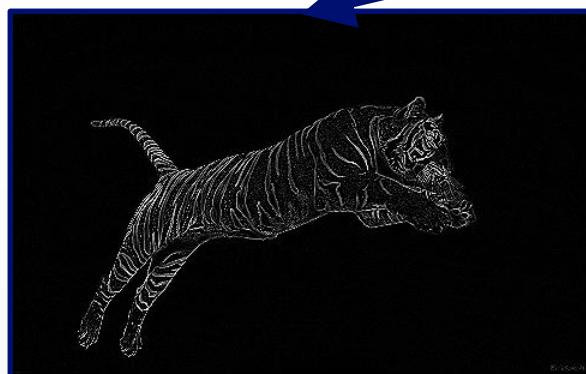
- Laplace operator:

$$\begin{bmatrix} 1 \\ -2 \\ 1 \end{bmatrix} + [1 \ -2 \ 1] = \boxed{\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}}$$

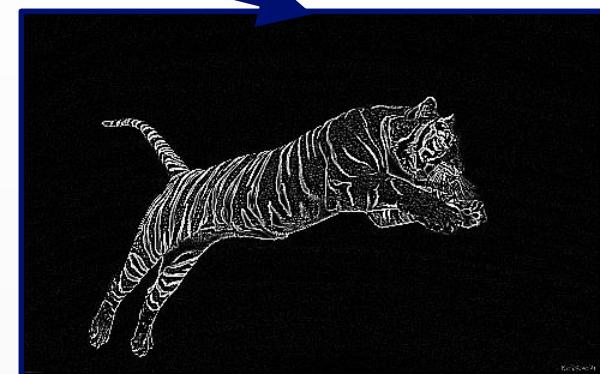
- There are other variations. (e.g. Second order Prewitt)

Edge Detection

- Edge detection with second order derivative:



Laplace edge detector

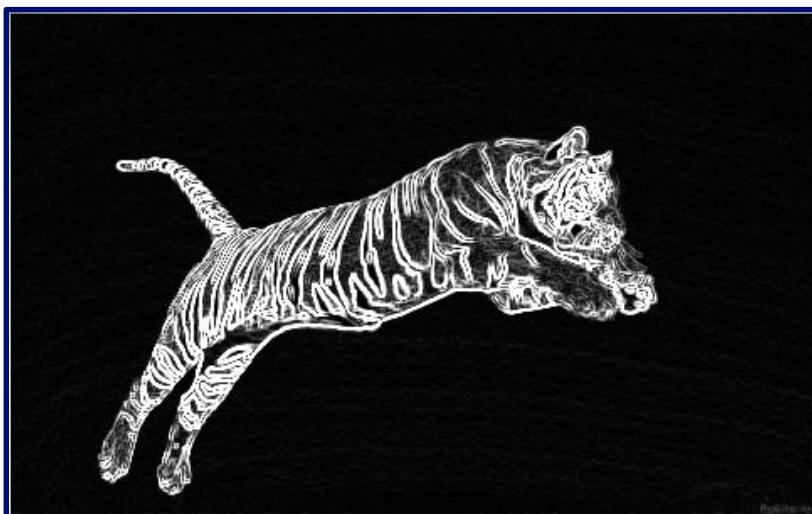


Prewitt 2nd order detector

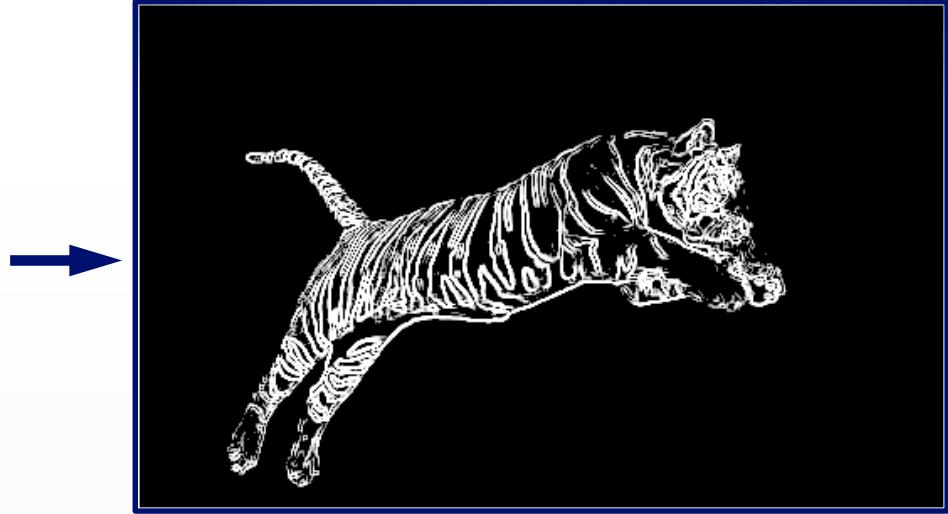
Edge Detection

◎ Thresholding:

- To eliminate weak edges, a threshold can be used on the gradient image:



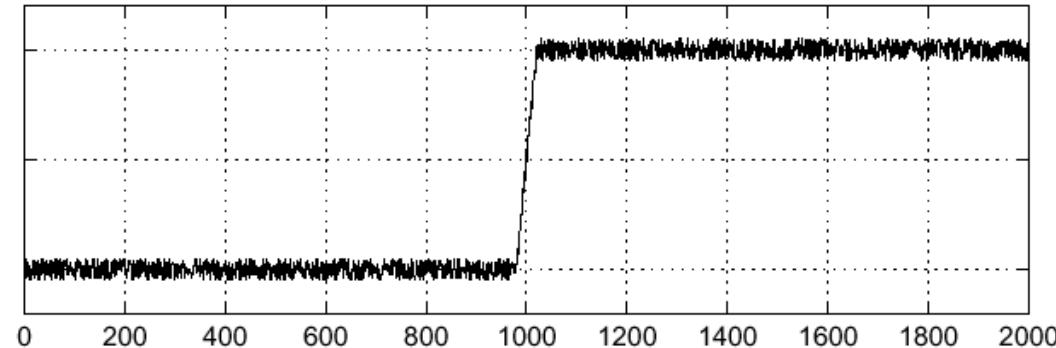
Prewitt first order gradient image



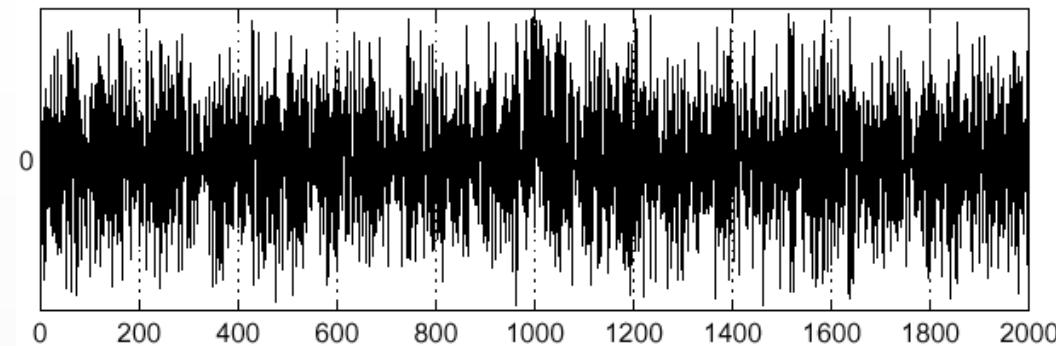
Prewitt first order gradient image
with threshold = 120

Noise filtering (1D demonstration)

$f(x)$



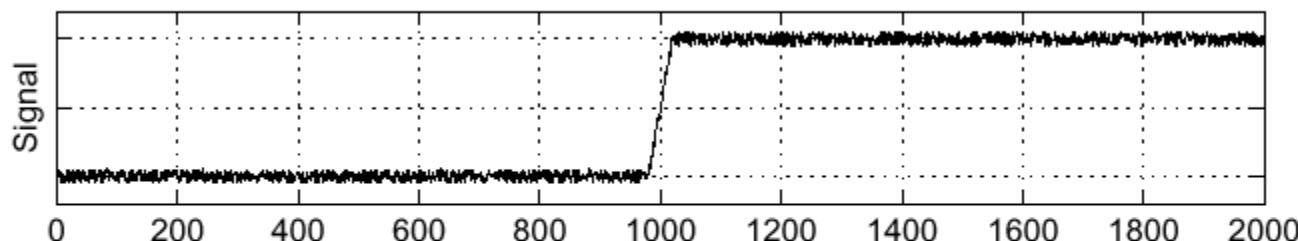
$\frac{d}{dx}f(x)$



- Where is the edge?

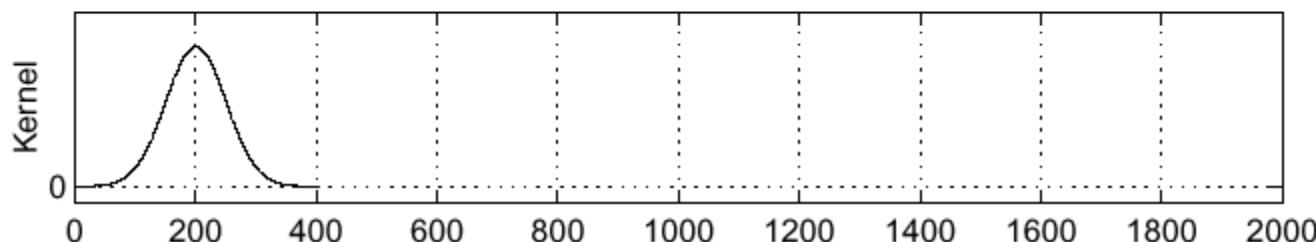
Sigma = 50

f



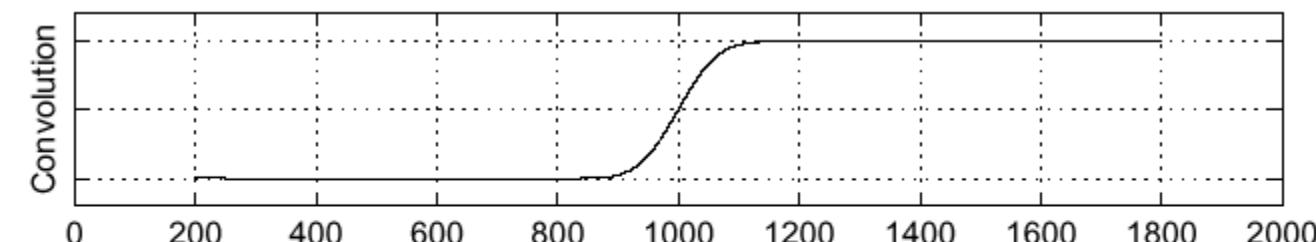
f: original signal

h



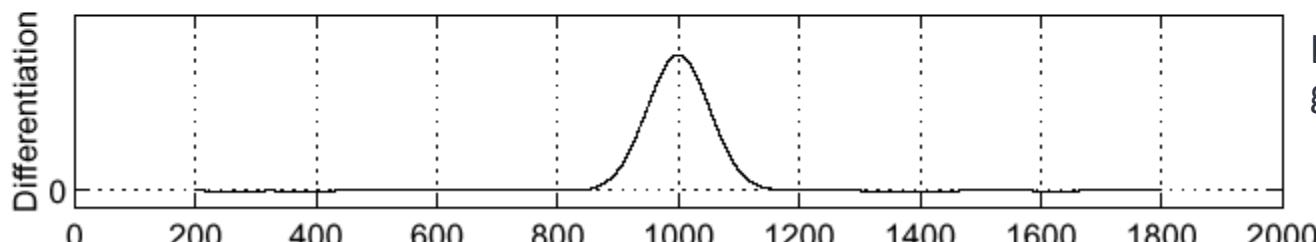
h: Gaussiam blur
kernel

h^*f



h^*f : filtered signal

$p^*(h^*f)$



p: Prewitt (first order
gradient kernel)

-1	0	1
----	---	---

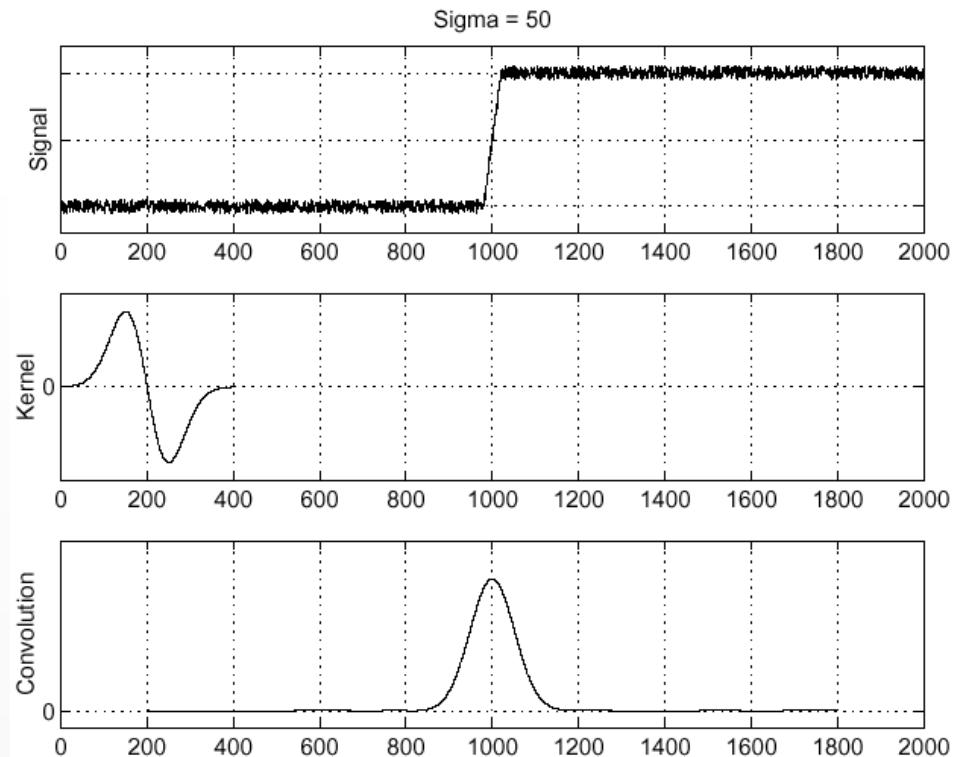
Smoothing the signal with Gaussian kernel, followed by applying a first order Prewitt kernel

Associativity of convolution: $p^*(h*f) = (p^*h)^*f$

$$\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x}h) * f$$

- No need for applying 2 convolutions, only one with the derivative of Gaussian operator (can also be approximated by a discrete kernel)

f
 $\frac{\partial}{\partial x}h \approx p^*h$
Derivative of the Gaussian kernel



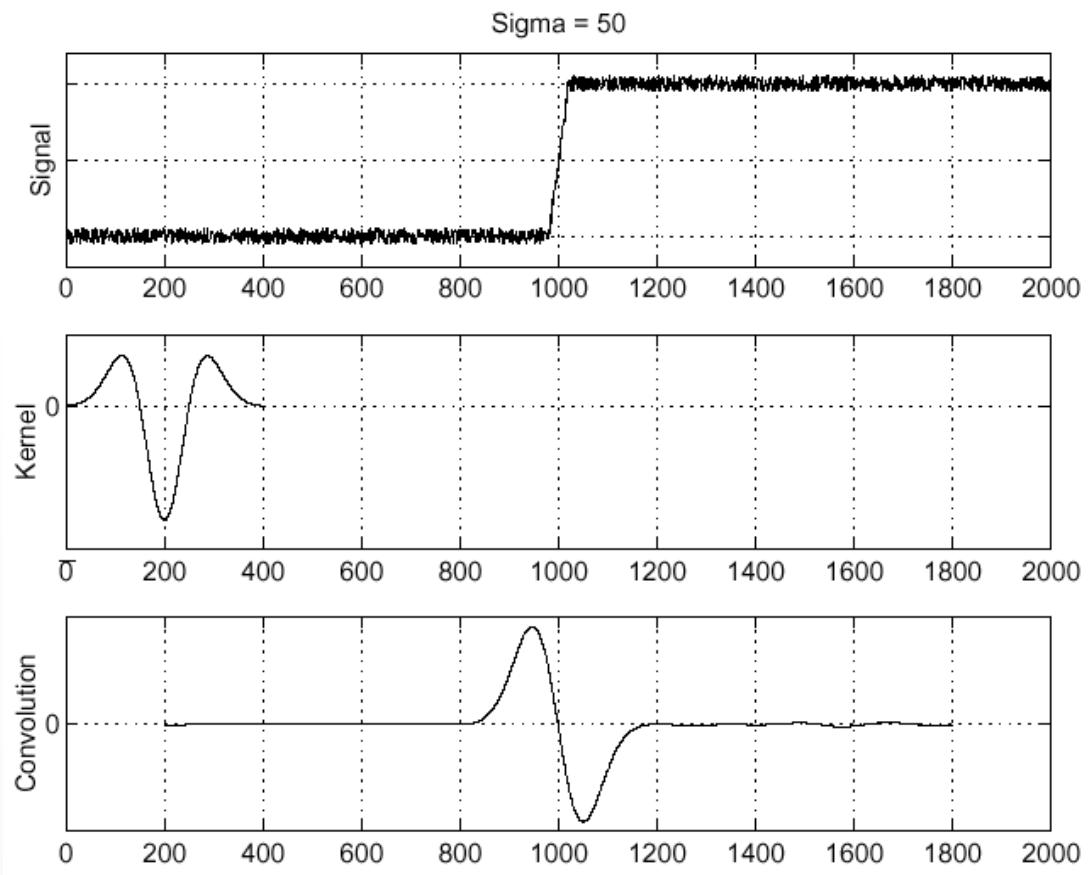
Second order case: Laplacian of Gaussian (LoG)

- Smoothing + Laplace = conv. with LoG operator

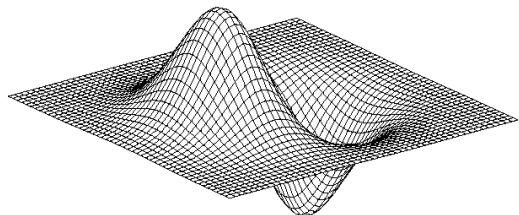
f
h: Gaussian smoothing kernel
l: Laplace-kernel

$$\frac{\partial^2}{\partial x^2} h \approx l * h$$

$$(\frac{\partial^2}{\partial x^2} h) \star f$$

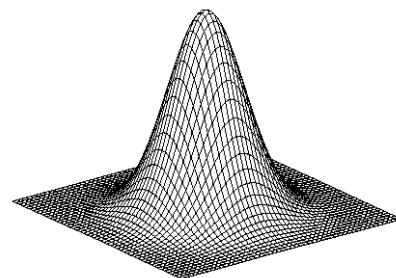


2D edge detection with filtering:



2D Gaussian smoothing kernel

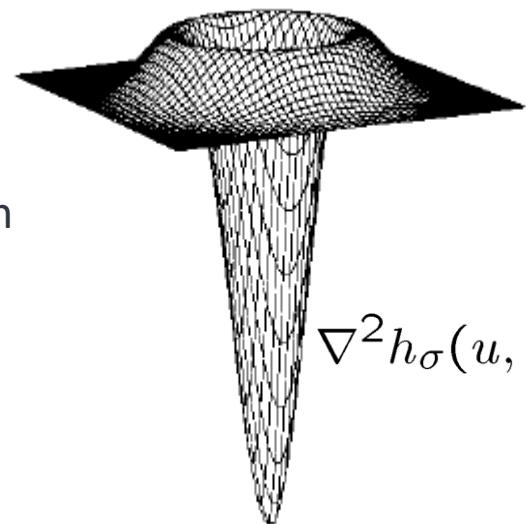
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



„Derivative“ of the 2D Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_\sigma(u, v)$$

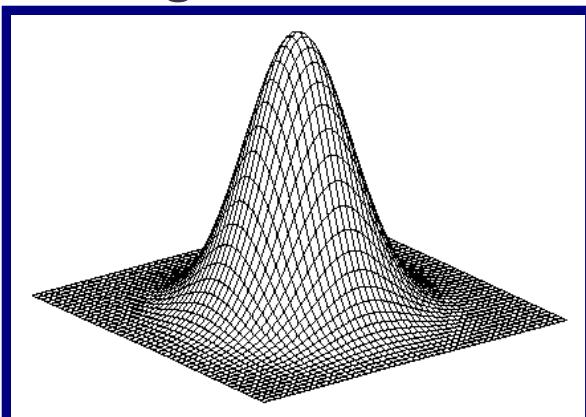
- ∇^2 henceforward the **Laplace** operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

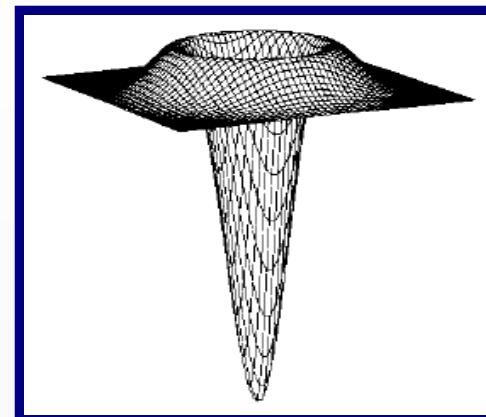
- Determining kernel coefficients with discrete approximation of the 2D function

Reducing the effect of noise on edge images

- Edge detection with noise reduction:
 - 1. step: Noise reduction by convolution with Gaussian filter
 - 2. step: Edge detection by convolution with Laplacian kernel
- Since convolution operation is associative we can convolve the Gaussian smoothing filter with the Laplacian filter first, and then convolve this hybrid filter (**Laplacian of Gaussian: LoG**) with the image.



Gaussian function

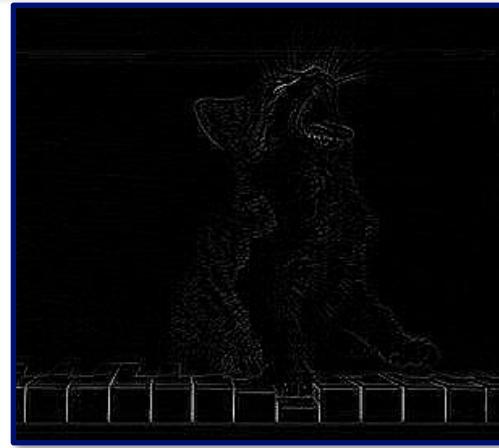


Laplacian of Gaussian

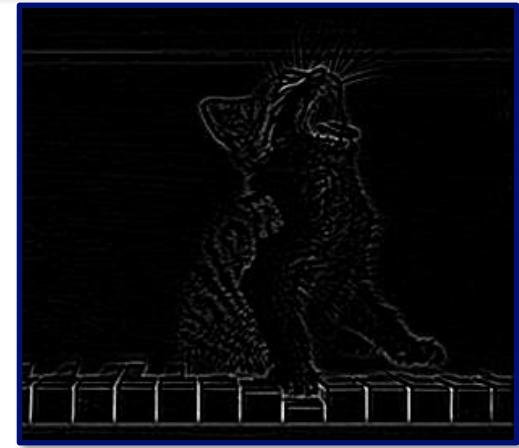
Laplacian of Gaussian



Original Image



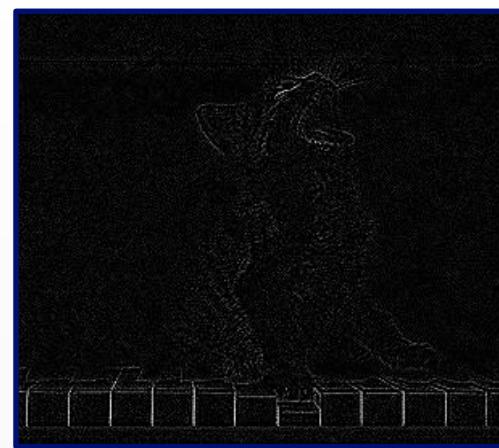
Laplacian



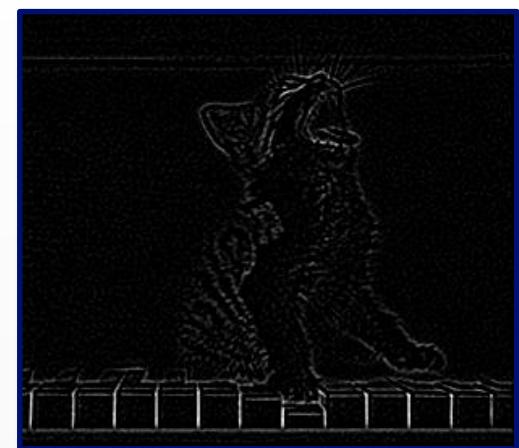
Laplacian of Gaussian



Noisy Image

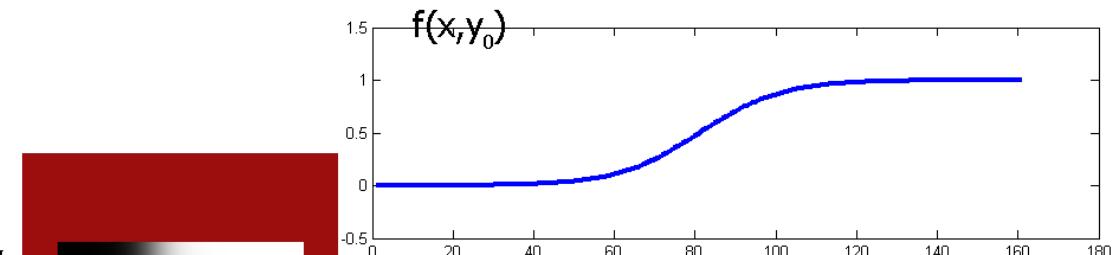


Laplacian



Laplacian of Gaussian

Edge crispening



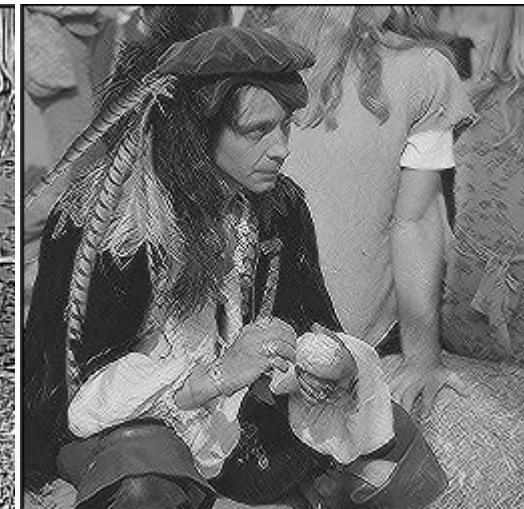
Convolution matrix:

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

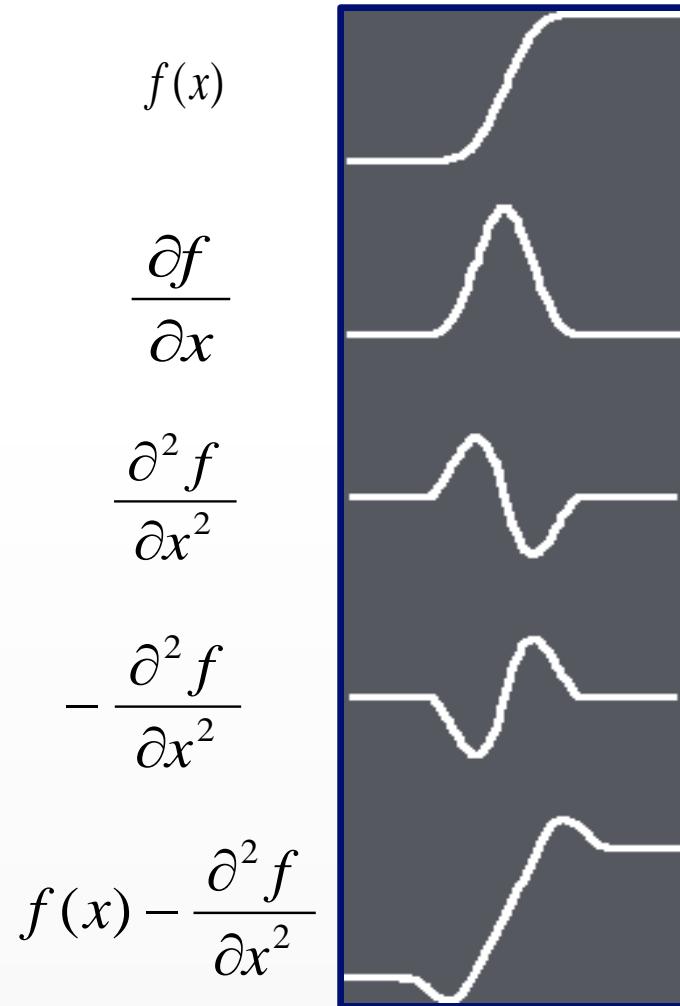
Edge crispening

- Often enhances the image quality

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix} \quad \begin{bmatrix} -1 & -1 & -1 \\ -1 & 9 & -1 \\ -1 & -1 & -1 \end{bmatrix} \quad \begin{bmatrix} 1 & -2 & 1 \\ -2 & 5 & -2 \\ 1 & -2 & 1 \end{bmatrix}$$



Edge Enhancement



Kernel for edge enhancement with Laplace operator:

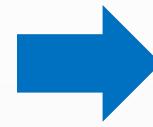
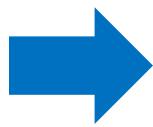
$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Original image



Edge enhanced image



Canny edge detector

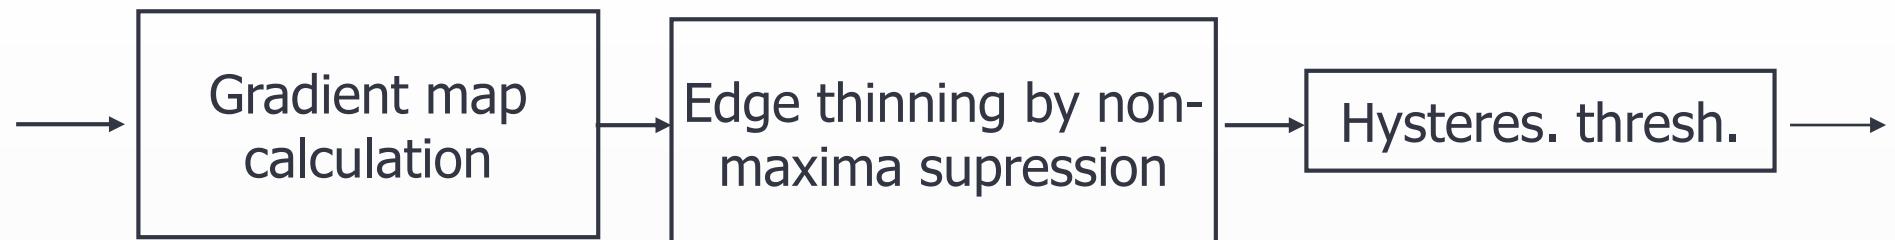
- Remember: properties of a good edge detector:

- Good detection:
 - detects as many real edges as possible
 - does not create false edges (because of e.g. image noise)
- Good localization:
 - the detected edges should be as close to the real edges as possible
- Isotropic:
 - all edges are detected regardless of their direction

- John F. **Canny** has developed an edge detector in 1986 to meet these requirements.

Canny edge detector

- Goal: extracting a **connected, *one-pixel-thick*** edge network
- Filtering Gaussian noise
- Three main steps:



Canny Edge detector

- >Main steps of the algorithm:

1. **Noise reduction:**

- The original image is convolved with a **Gaussian kernel** to reduce image noise.

2. **Gradient intensity and direction calculation:**

- The horizontal and vertical derivative image is calculated (e.g. with Prewitt kernel)
- Based on them the gradient intensity and direction can be calculated:

$$d^x = \left(\frac{\partial f}{\partial x} \right) \quad d = \|\nabla f\| = \sqrt{(d^x)^2 + (d^y)^2}$$

$$d^y = \left(\frac{\partial f}{\partial y} \right) \quad \Theta = \arctan \left(\frac{d^x}{d^y} \right)$$

Canny - 1st step: gradient map

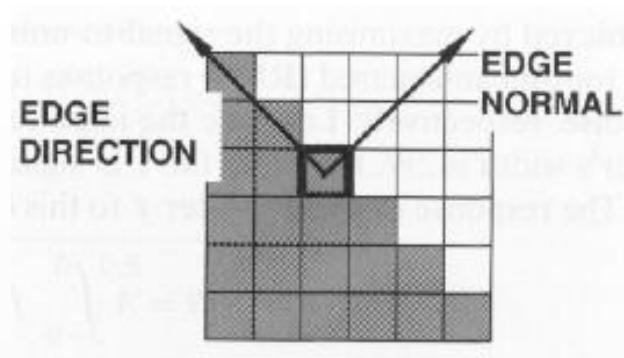
- Convolution (similarly as before)
 - Image smoothing with a Gaussian filter
 - x and y direction gradient estimation
- At each pixel (i,j) calculate gradient magnitude $[d(i,j)]$ and direction $[n(i,j) - \text{edge normal}]$

$$\|\nabla f\|_{ij} \propto d(i, j) = \sqrt{[d^x(i, j)]^2 + [d^y(i, j)]^2}$$

$$n(i, j) = \arctan\left(\frac{d^x(i, j)}{d^y(i, j)}\right)$$

Canny - 1st step: gradient map

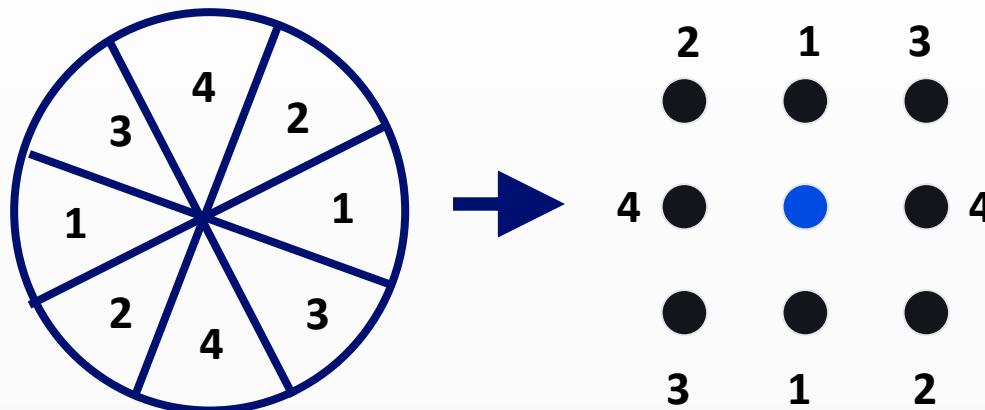
- ◎ $d(i,j)$ = edge magnitude (how sharp is the edge – proportional to the gradient magnitude)
- ◎ $n(i,j)$ = edge normal (which direction)



Canny Edge detector

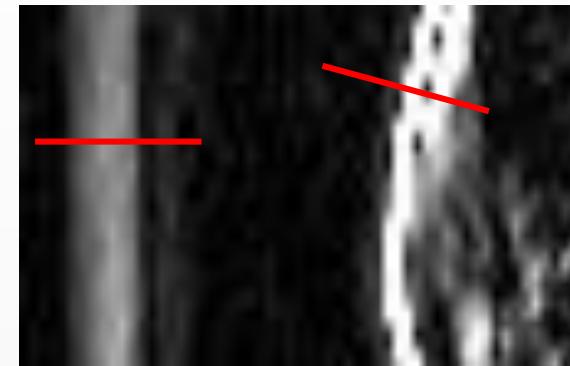
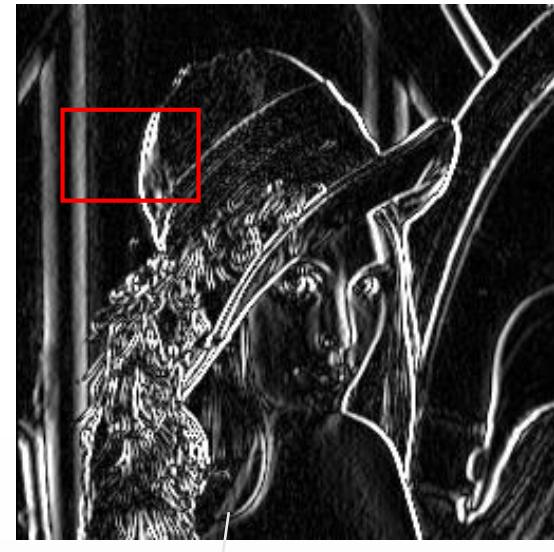
3. Non-Maximum Suppression:

- The goal is thinning the edges.
- Each edge is categorized into one of *4 main edge directions* (0° , 45° , 90° , 135°), based on the gradient direction image (θ).
- At every pixel, it suppresses the edge, by setting its value to 0, if its magnitude is not greater than the magnitude of the two neighbors in the gradient direction:



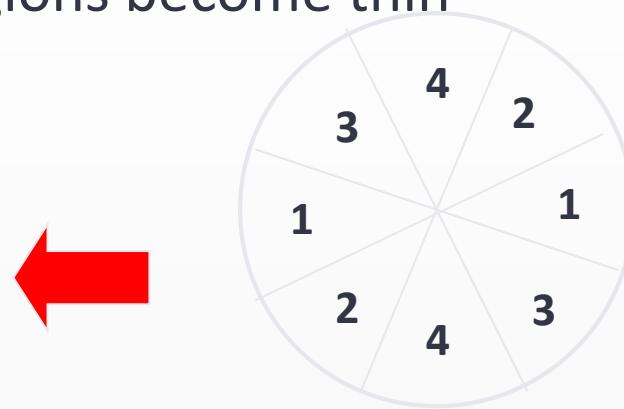
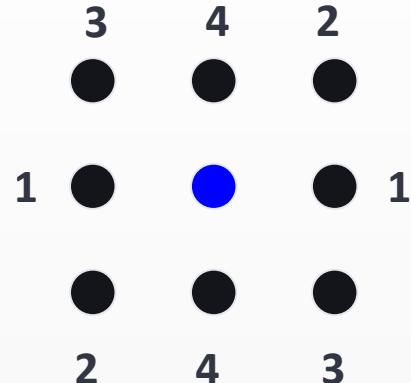
Canny – 2nd step: Non-max Suppression

- Goal: thinning the edges
- The gradient map may contain „thick” with large gradient values.
Earlier methods may classify all of these points as edges.
- Along the highlighted line segments perpendicular to the edges (marked with red) we should only mark a single point as edge point, the one which is locally the brightest



Canny – 2nd step: Non-max Suppression

1. Consider 4 principal directions: $(0^\circ, 45^\circ, 90^\circ, 135^\circ)$
2. To each pixel (i, j) we assign the principal direction $a(i, j)$, which one is the closest to local edge normal $n(i, j)$
3. If local edge magnitude $d(i, j)$ is smaller than in any neighboring pixel in the $a(i, j)$ set $G(i, j) := 0$. Otherwise (local max) : $G(i, j) := d(i, j)$
4. Result: G image obtained from the d gradient-magnitude map, where the edge-candidate regions become thin



Canny Edge detector

4. *Hysteresis thresholding:*

- Problem with simple thresholding:
 - if the threshold is low, many false edges will appear
 - if the threshold is high, true edges will disappear
- Solution: using two threshold instead of only one: t_1, t_2 , where $t_1 > t_2$
 - if the edge magnitude at (i, j) point is higher than t_1 then it is an edge
 - if the edge magnitude at (i, j) point is lower than t_2 then it is not an edge
 - if the edge magnitude at (i, j) point is lower than t_1 but higher than t_2 , then it is an edge, only if one of its neighbors in the direction of $\theta(i, j)$ is an edge.



original image



gradient intensity



non-maximum suppression



hysteresis thresholding

Canny – 3rd step: Thresholding

- ◎ Naive solution: thresholding the G map with a threshold t
 - If t is too small, we obtain many false edge points. If t is too large: valid edges disappear.
 - If the gradient magnitude of the edges fluctuates around the threshold, many disruptions (broken edge segments) may appear
- ◎ Improved solution: hysteresis thresholding
 - Using 2 thresholds t_1 and t_2 ($t_1 < t_2$):
 - If the value of $G(x, y)$ is larger than t_2 , (x, y) is certainly edge point
 - If the value of $G(x, y)$ is smaller than t_1 , (x, y) is certainly not an edge point
 - If the value of $G(x, y)$ is between the threshold, we mark it as edge point if and only if it has a neighboring pixel already classified as edge in the direction perpendicular to the **edge normal**

Canny edge detector - result



Input image

Canny edge detector - result



Norm of gradient: „d”

Canny edge detector - result



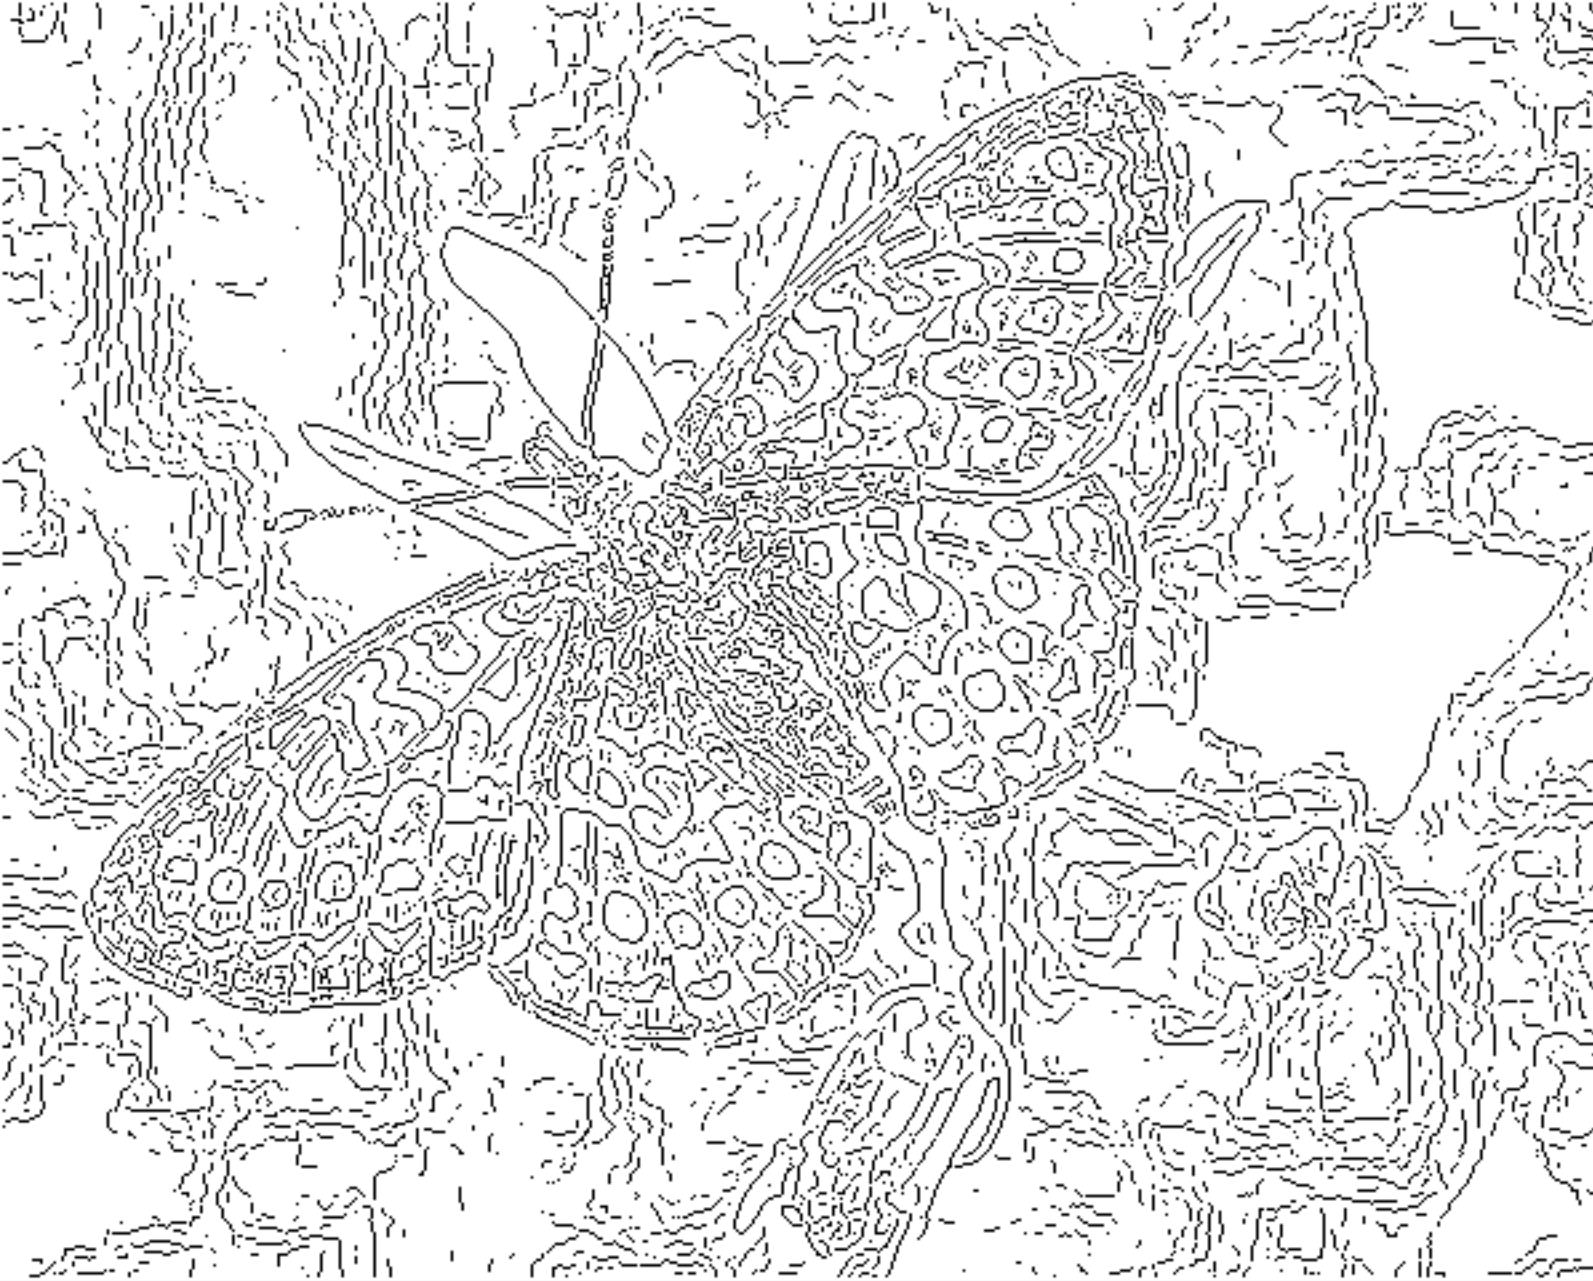
After thinning (E) (non-maximum suppression)

Canny edge detector - result



hysteresis thresholding





fine scale
high
threshold



A high-resolution grayscale image of a human brain's cortical surface. The image shows intricate folds and grooves (gyri and sulci) in great detail. The brain is oriented with the left hemisphere on the left and the right hemisphere on the right. The overall texture is highly detailed, with numerous small circular features representing individual blood vessels or capillaries.

coarse
scale,
high
threshold

Canny edge detector - results

