

A Fast, Vectorizable Algorithm for Producing Single-Precision Sine-Cosine Pairs

Marcus H. Mendenhall

Abstract—This paper presents an algorithm for computing Sine-Cosine pairs to modest accuracy, but in a manner which contains no conditional tests or branching, making it highly amenable to vectorization. An exemplary implementation for PowerPC AltiVec processors is included, but the algorithm should be easily portable to other architectures, such as Intel SSE.

Index Terms—Elementary function approximation, Numerical algorithms, Parallel Algorithms, Mathematical Software / Parallel and vector implementations

I. INTRODUCTION

THE need for sine-cosine pairs arises often in numerical computing and digital signal processing. For many of the applications, the functions are needed at uniformly spaced values of the argument, and can be generated by efficient and accurate recursion relations.

For some applications, though, one needs to compute trigonometric functions in random order, or for non-uniformly spaced values. If a large number of such evaluations are needed, it would be convenient to be able to take advantage of the extremely fast vector processing units included on many modern desktop computer families. Unfortunately, most common algorithms require that one reduce the angle modulo $\pi/2$ onto the range $[-\pi/4, \pi/4)$ (or sometimes mod π onto $[-\pi/2, \pi/2)$) [1], then to compute the function using some reasonable power series, and then to fill in the rest of the circle by appropriate exchanges and sign changes. The shuffling at the end to fill in all four quadrants is not easy to do if multiple angles are being computed at once in a vector processor, since each quadrant requires a different shuffle.

Marcus Mendenhall is with the Vanderbilt University W. M. Keck Free Electron Laser Center and the Department of Physics and Astronomy. Research Supported by the United States DOD MFEL program under grant FA9550-04-1-0045

II. METHODOLOGY

The approach taken in this work is slightly different than the approach used historically. The goal of it is to avoid the relatively expensive testing and shuffling resulting from the reduction of the angle to one quadrant. To accomplish this, the angle is first scaled by $1/2\pi$ and reduced modulo 1 onto $[-1/2, 1/2)$ by rounding and subtraction. Then, a power series of the functions of an angle scaled down by a power of two is computed, and the resulting trig pair is shifted back onto the whole circle using double-angle formula recursion.

Computing functions in this manner has a number of useful properties. As discussed, it is highly vectorizable. It also produces a very smooth result; that is, except for the error when the modular conversion wraps around at $(2n + 1)\pi$, there are no seams at which one must eliminate discontinuities. Thus, algorithms which depend on differentiation (e.g. nonlinear least-squares fitting) will not be likely to be disrupted by the joining. In reality, most IEEE-754 compliant algorithms put a great deal of effort into avoiding such discontinuities by assuring the error in the computation is less than 0.5 LSB. This effort, though costs time. The method described here, instead, sacrifices a little absolute accuracy in trade for extremely high speed, while preserving smoothness.

In the implementation shown in this paper, I compute $\sin(2\pi x/4)$ and $\cos(2\pi x/4)$ where x is the scaled and reduced angle as discussed above, and then use 2 angle-doublings to expand this over the entire circle using:

$$\begin{aligned}\sin(2\theta) &= 2 \sin \theta \cos \theta \\ \cos(2\theta) &= \cos^2 \theta - \sin^2 \theta\end{aligned}\tag{1}$$

The issue that needs to be addressed when carrying out any recursion relation such as the repeated double-angle relations above is that of stability.

Since this algorithm is designed to generate single-precision functions, and to operate on a single-precision vector processing unit, one must pay close attention to instabilities in this recursion due to roundoff error. In fact, for this algorithm, there were two candidates for the angle doubling. The first is the one above, and the second is

$$\begin{aligned}\sin(2\theta) &= 2 \sin \theta \cos \theta \\ \cos(2\theta) &= 1 - 2 \sin^2 \theta\end{aligned}\quad (2)$$

Each of the two methods has an instability, but in the chosen method it is easily fixed at the end, and in the rejected method it is not so.

Consider the case of a small error introduced: define

$$s_1 = a \sin(\theta + \delta), \quad c_1 = a \cos(\theta + \delta)$$

and iterate once using (1) to get:

$$\begin{aligned}s_2 &= 2a^2 \cos(\theta + \delta) \sin(\theta + \delta) = a^2 \sin 2(\theta + \delta) \\ c_2 &= a^2 \cos^2(\theta + \delta) - a^2 \sin^2(\theta + \delta) \\ &= a^2 \cos 2(\theta + \delta)\end{aligned}$$

Note that, then, the magnitude has drifted, but the error in the angle is stable.

Now, consider the case using (2):

$$\begin{aligned}s_2 &= a^2 \sin 2(\theta + \delta) \\ c_2 &= 1 - 2a^2 \sin^2(\theta + \delta) \\ &= a^2 (\cos^2(\theta + \delta) - \sin^2(\theta + \delta)) + (1 - a^2) \\ &= a^2 \cos 2(\theta + \delta) + (1 - a^2)\end{aligned}$$

thus mixing the angle and amplitude error inextricably.

Thus, the form of (1) produces a pair (after n iterations) of $a^{2n} \cos 2^n(\theta + \delta)$, $a^{2n} \sin 2^n(\theta + \delta)$ which has the same relative error in the angle as the initial estimate, but a scale factor. Since the scale factor has no effect on the accuracy of the properly-normalized trig pair, there are two options for computing the raw pair. These options differ in the fitting algorithm used to find the coefficients for the series.

The first method tries to produce a pair which has a very small normalization error by directly computing least-squares series for sine and cosine. Assuming the original amplitude error is small, so $a = 1 + \epsilon$, then $a^{2n} \approx 1 + 2n\epsilon$. Thus, the

final amplitude has a relative error $2n$ times bigger than the initial error. For $n = 2$ and single-precision arithmetic, this implies an amplitude error of about 4×10^{-7} (assuming 1 LSB error at the start). To correct for this at the end, compute $s_n^2 + c_n^2 = 1 + \alpha$ and divide the function pair by $\sqrt{1 + \alpha}$. In practice, since α is small, multiplying the pair by $1 - \alpha/2 = (3 - s^2 - c^2)/2$ provides a correctly normalized result. Using this method, the RMS error is 1.2×10^{-7} and the maximum error is 4.8×10^{-7} . These errors are measured as $\sqrt{(\cos \theta - \text{ref_} \cos \theta)^2 + (\sin \theta - \text{ref_} \sin \theta)^2}$ where the reference functions are computed to double precision using the standard system library calls. The maximum amplitude error, computed as $1 - \sqrt{\cos^2 \theta + \sin^2 \theta}$ is 1.8×10^{-7} . For this method, the polynomials:

$$\begin{aligned}\sin \theta/4 &= x(1.5707963235 \\ &\quad - 0.645963615 x^2 \\ &\quad + 0.0796819754 x^4 \\ &\quad - 0.0046075748 x^6) \\ \cos \theta/4 &= 1 \\ &\quad - 1.2336977925 x^2 \\ &\quad + 0.2536086171 x^4 \\ &\quad - 0.0204391631 x^6\end{aligned}$$

are used, where x is the range-reduced value of $\theta/2\pi$ discussed at the beginning of this paper.

The second method relaxes the requirement that the magnitude error be small, by fitting the input angle to $\arctan(\sin \theta / \cos \theta)$, but requires then an accurate division by the actual magnitude of the pair at the end. Doing this produces a slightly better error budget, with an RMS error of 9.8×10^{-8} and a maximum error of 3.8×10^{-7} . The penalty in the final algorithm is doing a precise reciprocal and multiply, which costs about 5% in speed. For this method, the polynomials:

$$\begin{aligned}\sin^\dagger \theta/4 &= x(1.5707963268 \\ &\quad - 0.6466386936 x^2 \\ &\quad + 0.0679105987 x^4 \\ &\quad - 0.0011573807 x^6) \\ \cos^\dagger \theta/4 &= 1 \\ &\quad - 1.2341299769 x^2 \\ &\quad + 0.2465220241 x^4 \\ &\quad - 0.0123926179 x^6\end{aligned}$$

are used. Note that calling these sine and cosine is somewhat dangerous; they are really the numerator and denominator of a good rational-function representation for $\tan \theta/4$ but do not satisfy $\sin^2 + \cos^2 = 1$. I have annotated the terms in the equations with the (\dagger) symbol to remind the reader of this.

III. RESULTS

In figure 1, I show a sample implementation of this, coded in C (compatible with gcc3.xx or IBM xlc). This code computes a trig pair vector (4 trig pairs) in 140 ns on a 500 MHz PowerPC 7400 (aka PowerMacintosh G4).

It is important to note that, even for this algorithm, the throughput can be improved quite a bit. Because of the chain-computation nature of this, there are many stalls waiting for latency in the vector unit. If one computes two vectors at a time, by just interleaving this code with itself with different variable names, throughput rises another 50%. However, the logistics of using this in other code often does not make it worthwhile. Also, the nature of optimizations of this type is quite architecture-dependent and therefore not within the intended scope of this work. However, with this optimization, one can generate a single-precision trig function every 7 clock cycles of the CPU (8 pairs in 100 ns), including loading and storing the results. For a discussion of the instruction set and throughput of this particular vector unit, see [2].

IV. CONCLUSIONS

For applications requiring very high speed, random-access calculations of sine-cosine pairs, it is possible to forego a very small amount of accuracy in exchange for a great deal of speed. Using multiple-angle formulas to expand a sine-cosine pair computed on a small piece of the reduced range is an effective approach to this. The author has been applying this method to the rapid computation of optical diffraction patterns via Huygen's principle, but it should be widely applicable to other problems for which very large numbers of single-precision trigonometric pairs are needed.

REFERENCES

- [1] for example: Abramowitz, M. and Stegun, I., "Pocketbook of Mathematical Functions", Verlag Harri Deutsch, 1984, eq. 4.3.97 and 4.3.98

- [2] "AltivecTM Technology Programming Interface Manual", Motorola SPS document ID ALTIVECPIM/D, 1999

PLACE
PHOTO
HERE

Marcus H. Mendenhall is a Research Associate Professor of Physics at Vanderbilt University. He received his Ph.D. from CalTech in 1983. He has been working recently on pulsed, tunable Xray sources based on Compton backscattering of laser photons from a relativistic electron beam. Techniques of numerical computation and analysis are a major side-interest.

```

/* define FASTER_SINCOS for the slightly-less-accurate results in slightly less time */
#define FASTER_SINCOS

#if !defined(FASTER_SINCOS) /* these coefficients generate a badly un-normalized sine-cosine pair, but the angle
#define ss1 1.5707963268
#define ss2 -0.6466386396
#define ss3 0.0679105987
#define ss4 -0.0011573807
#define cc1 -1.2341299769
#define cc2 0.2465220241
#define cc3 -0.0123926179
#else /* use 20031003 coefficients for fast, normalized series*/
#define ss1 1.5707963235
#define ss2 -0.645963615
#define ss3 0.0796819754
#define ss4 -0.0046075748
#define cc1 -1.2336977925
#define cc2 0.2536086171
#define cc3 -0.0204391631
#endif
inline void FastSinCos(vector float v, struct phase *ph)
{
    vector float s1, s2, c1, c2, fixmag1;

    vector float x1=vec_madd(v, (vector float)(1.0/(2.0*3.1415926536)),
        (vector float)(0.0));

    /* q1=x/2pi reduced onto (-0.5,0.5), q2=q1**2 */
    vector float q1=vec_nmsub(vec_round(x1), (vector float)(1.0), x1);
    vector float q2=vec_madd(q1, q1, (vector float)(0.0));

    s1=    vec_madd(q1,
        vec_madd(q2,
            vec_madd(q2,
                vec_madd(q2, (vector float)(ss4),
                    (vector float)(ss3)),
                (vector float)( ss2)),
            (vector float)(ss1)),
        (vector float)(0.0));
    c1=    vec_madd(q2,
        vec_madd(q2,
            vec_madd(q2, (vector float)(cc3),
                (vector float)(cc2)),
            (vector float)(cc1)),
        (vector float)(1.0));
    /* now, do one out of two angle-doublings to get sin & cos theta/2 */
    c2=vec_nmsub(s1, s1, vec_madd(c1, c1, (vector float)(0.0)));
    s2=vec_madd((vector float)(2.0), vec_madd(s1, c1, (vector float)(0.0)),
        (vector float)(0.0));
    /* now, cheat on the correction for magnitude drift...
       if the pair has drifted to (1+e)*(cos, sin),
       the next iteration will be (1+e)**2*(cos, sin)
       which is, for small e, (1+2e)*(cos,sin).
       However, on the (1+e) error iteration,
       sin**2+cos**2=(1+e)**2=1+2e also,
       so the error in the square of this term
       will be exactly the error in the magnitude of the next term.
       Then, multiply final result by (1-e) to correct */

    #if defined(FASTER_SINCOS) /* this works with properly normalized sine-cosine functions, but un-normalized is mor
        fixmag1=vec_nmsub(s2,s2, vec_nmsub(c2, c2, (vector float)(2.0)));
    #else /* must use this method with un-normalized series, since magnitude error is large */
        fixmag1=Reciprocal(vec_madd(s2,s2,vec_madd(c2,c2,(vector float)(0.0))));
    #endif

    c1=vec_nmsub(s2, s2, vec_madd(c2, c2, (vector float)(0.0)));
    s1=vec_madd((vector float)(2.0), vec_madd(s2, c2, (vector float)(0.0)),
        (vector float)(0.0));

    ph->c=vec_madd(c1, fixmag1, (vector float)(0.0));
    ph->s=vec_madd(s1, fixmag1, (vector float)(0.0));
}

```

Fig. 1. Sample Implementation for PowerPC AltiVec vector processor